Specware 4.2 Language Manual

Specware 4.2 Language Manual

Copyright © 2001-2007 by Kestrel Development Corporation
Copyright © 2001-2007 by Kestrel Technology LLC

The name Specware® is a registered trademark of Kestrel Development Corporation

Table of Contents

Disclaimer

As experience is gained with Specware 4.2, both the operation of the Specware system and the Metaslang language are bound to undergo changes, which may not always be fully "backwards compatible".

For updates, news and bug reports, visit the Specware web site http://www.specware.org.

Disclaimer

Chapter 1. Introduction to Specware

1.1. What Is Specware?

Specware is a tool for building and manipulating a collection of related specifications. Specware can be considered:

- a design tool, because it can represent and manipulate designs for complex systems, software or otherwise

- a logic, because it can describe concepts in a formal language with rules of deduction

- a programming language, because it can express programs and their properties

- a database, because it can store and manipulate collections of concepts, facts, and relationships

Specifications are the primary units of information in Specware. A specification, or theory, describes a concept to some degree of detail. To add properties and extend definitions, you create new specifications that import or combine earlier specifications. Within a specification, you can reason about objects and their relationships. You declare types (data types) and operations (ops, functions), axioms that state properties of operations, and theorems that follow logically from axioms.

A morphism is a relationship between specifications that describes how the properties of one map relate to the properties of another. Morphisms describe both part-of and is-a relationships. You can propagate theorems from one specification to another using morphisms; for example, if the QEII is a ship, and ships cannot fly, then the QEII cannot fly.

1.2. What Is Specware For?

Specware is a general-purpose tool that you can use to develop specifications for any system or realm of knowledge. You can do this as an abstract process, with no reference to computer programming; or you can produce a computer program that is provably a correct implementation of a specification; or you can use the process to redesign an existing program.

You can use Specware to:

- **Develop domain theories**

  You can use Specware to do "ontological engineering" -- that is, to describe a real-world domain of knowledge in explicit or rigorous terms. You might wish to develop a domain theory in abstract terms that are not necessarily intended to become a computer program. You can use the inference engine to test the internal logic of your theory, derive conclusions, and propose theorems.

  You can use specifications and morphisms to represent abstract knowledge, with no refinement to any kind of concrete implementation.

  More commonly, you would use Specware to model expert knowledge of engineering design. In this case you would refine your theoretical specifications and morphisms to more concrete levels.

- **Develop code from specifications**

  You can use Specware to develop computer programs from specifications. One advantage of using Specware for this task is that you can prove that the generated code does implement the specification correctly. Another advantage is that you can develop and compare different implementations of the same specification.

- **Develop specifications from code**

  You can use Specware for reverse engineering -- that is, to help you derive a specification from existing code. To do this, you must examine the code to determine what problems are being solved by it, then use Specware's language Metaslang to express the problems as specifications. In addition to providing a notation tool for this process, Specware allows you to operate on the derived specification. Once you have derived a specification from the original code, you can analyze the specification for correctness and completeness, and also generate different and correct implementations for it.

1.3. The Design Process in Specware

To solve real problems, programs typically combine domain theories about the physical world with problem solving theories about the computational world. Your domain theory is an abstract representation of a real-world problem domain. To implement it,

you must transform the domain theory to a concrete computational model. The built-in specification libraries describe mathematical and computational concepts, which are building blocks for an implementation. Your specifications combine real-world knowledge with this built-in computational knowledge to generate program code that solves real-world problems in a rigorous and provable way.

You interpret designs relative to an initial universe of models. In software design, for example, the models are programs, while in engineering design, they are circuits or pieces of metal. To design an object is to choose it from among the universe of possible models. You make this choice by beginning with an initial description and augmenting it until it uniquely describes the model you desire. In Specware, this process is called refinement.

Composition and refinement are the basic techniques of application building in Specware. You compose simpler specifications into more complex ones, and refine more abstract specifications into more concrete ones. When you refine a specification, you create a more specific case of it; that is, you reduce the number of possible models of it.

The process of refinement is also one of composition. To begin the refinement, you construct primitive refinements that show how to implement an abstract concept in terms of a concrete concept. You then compose refinements to deepen and widen the refinement.

For example, suppose you are designing a house. A wide but not deep view of the design specifies several rooms but gives no details. A deep but not wide view of the design specifies one room in complete detail. To complete the refinement, you must create a view that is both wide and deep; however, it makes no difference which view you create first.

The final refinement implements a complex, abstract specification from which code can be generated.

1.4. Stages of Application Building

Conceptually, there are two major stages in producing a Specware application. In the actual process, steps from these two stages may alternate.

1. Building a specification

2. Refining your specifications to constructive specifications

1.4.1. Building a Specification

You must build a specification that describes your domain theory in rigorous terms. To do this, you first create small specifications for basic, abstract concepts, then specialize and combine these to make them more concrete and complex.

To relate concepts to each other in Specware, you use specification morphisms. A specification morphism shows how one concept is a specialization or part of another. For example, the concept "fast car" specializes both "car" and "fast thing". The concept "room" is part of the concept "house". You can specialize "room" in different ways, one for each room of the house.

You specialize in order to derive a more concrete specification from a more abstract specification. Because the specialization relation is transitive (if A specializes B and B specializes C, then A specializes C as well), you can combine a series of morphisms to achieve a step-wise refinement of abstract specifications into increasingly concrete ones.

You combine specifications in order to construct a more complex concept from a collection of simpler parts. In general, you increase the complexity of a specification by adding more structural detail.

Specware helps you to handle complexity and scale by providing composition operators that take small specifications and combine them in a rigorous way to produce a complex specification that retains the logic of its parts. Specware provides several techniques for combining specifications, that can be used in combination:

- The import operation allows you to include earlier specifications in a later one.

- The translate operation allows you to rename the parts of a specification.

- The colimit operation glues concepts together into a shared union along shared subconcepts.

A shared union specification combines specializations of a concept. For example, if you combine "red car" with "fast car" sharing the concept "car", you obtain the shared union "fast, red car". If you combine them sharing nothing, you obtain "red car and fast car", which is two cars rather than one. Both choices are available.

1.4.2. Refining Your Specifications to Constructive Specifications

You combine specifications to extend the refinement iteratively. The goal is to create a

refinement between the abstract specification of your problem domain and a concrete implementation of the problem solution in terms of types and operations that ultimately are defined in the Specware libraries of mathematical and computational theories.

For example, suppose you want to create a specification for a card game. An abstract specification of a card game would include concepts such as card, suit, and hand. A refinement for this specification might map cards to natural numbers and hands to lists containing natural numbers.

The Specware libraries contains constructive specifications for various types, including natural numbers and lists.

To refine your abstract specification, you build a refinement between the abstract Hand specification and the List-based specification. When all types and operations are refined to concrete library-defined types and operations, the Specware system can generate code from the specification.

1.5. Reasoning About Your Code

Writing software in Metaslang, the specification and programming language used in Specware, brings many advantages. Along with the previously-mentioned possibilities for incremental development, you have the option to perform rigorous verification of the design and implementation of your code, leading to the a high level of assurance in the correctness of the final application.

1.5.1. Abstractness in Specware

Specware allows you to work directly with abstract concepts independently of implementation decisions. You can put off making implementation decisions by describing the problem domain in general terms, specifying only those properties you need for the task at hand.

In most languages, you can declare either everything about a function or nothing about it. That is, you can declare only its type, or its complete definition. In Specware you must declare the signature of an operation, but after that you have almost complete freedom in stating properties of the operation. You can declare nothing or anything about it. Any properties you have stated can be used for program transformation.

For example, you can describe how squaring distributes over multiplication:

```
axiom square_distributes_over_times is
  fa(a, b) square(a * b) = square(a) * square(b)
```

This property is not a complete definition of the squaring operation, but it is true. The truth of this axiom must be preserved as you refine the operation. However, unless you are going to generate code for `square`, you do not need to supply a complete definition.

The completeness of your definitions determines the extent to which you can generate code. A complete refinement must completely define the operations of the source theory in terms of the target theory. This guarantees that, if the target is implementable, the source is also implementable.

1.5.2. Logical Inference in Specware

Specware performs inference using external theorem provers; the distribution includes SRI's SNARK theorem prover. External provers are connected to Specware through logic morphisms, which relate logics to each other.

You can apply external reasoning agents to refinements in different ways (although only verification is fully implemented in the current release of Specware).

- Verification tests the correctness of a refinement. For example, you can prove that quicksort is a correct refinement of the sorting specification.

- Simplification is a complexity-reducing refinement. For example, given appropriate axioms, you can rewrite `3*a+3*b` to `3*(a+b)`.

- Synthesis derives a refinement for a given specification by combining the domain theory with computational theory. For example, you can derive quicksort semi-automatically from the sorting specification as a solution to a sorting problem, if you describe exactly how the problem is a sorting problem.

Chapter 2. Metaslang

This chapter introduces the Metaslang specification language.

The following sections give the grammar rules and meaning for each Metaslang language construct.

2.1. Preliminaries

2.1.1. The Grammar Description Formalism

The grammar rules used to describe the Metaslang language use the conventions of (extended) BNF. For example, a grammar rule like:

> wiffle ::= waffle [waffle-tail] | piffle { + piffle }*

defines a **wiffle** to be: either a **waffle** optionally followed by a **waffle-tail**, or a sequence of one or more **piffles** separated by terminal symbols + . (Further rules would be needed for **waffle, waffle-tail** and **piffle**.) In a grammar rule the left-hand side of ::= shows the kind of construct being defined, and the right-hand side shows how it is defined in terms of other constructs. The sign | separates alternatives, the square brackets [...] enclose optional parts, and the curly braces plus asterisk { ... }* enclose a part that may be repeated any number of times, including zero times. All other signs stand for themselves, like the symbol + in the example rule above.

In the grammar rules terminal symbols appear in a bold font. Some of the terminal symbols used, like | and {, are very similar to the grammar signs like | and { as described above. They can hopefully be distinguished by their bold appearance.

Grammar rules may be *recursive*: a construct may be defined in terms of itself, directly or indirectly. For example, given the rule:

> piffle ::= **1** | **M** { piffle }*

here are some possible **piffles**:

> 1   M   M1   M111   MMMM   M1M1

Note that the last two examples of **piffles** are ambiguous. For example, M1M1 can be interpreted as: M followed by the two **piffles** 1 and M1, but also as: M followed by the

three piffles 1, M, and another 1. Some of the actual grammar rules allow ambiguities; the accompanying text will indicate how they are to be resolved.

2.1.2. Models

```
op ::= op-name
```

In Metaslang, *op* is used used as an abbreviation for "op-name", where op-names are declared names representing values. (*Op* for *operator*, a term used for historical reasons, although including things normally not considered operators.)

```
spec ::= spec-form
```

The term spec is used as short for spec-form. The *semantics* of Metaslang specs is given in terms of classes of *models*. A model is an assignment of sets of values (called "types") to all the type-names and of "typed" values to all the ops declared -- explicitly or implicitly -- in the spec. The notion of *value* includes numbers, strings, arrays, functions, etcetera. A typed value can be thought of as a pair (T, V), in which T is a type and V is a value that is an inhabitant of T. For example, the expressions 0 : Nat and 0 : Integer correspond, semantically, to the typed values $(N, 0)$ and $(Z, 0)$, respectively, in which N stands for the set of natural numbers $\{0, 1, 2, ...\}$, and Z for the set of integers $\{..., -2, -1, 0, 1, 2, ...\}$. For example, given this spec:

```
spec
  type Even
  op next : Even -> Even
  axiom nextEffect is
    fa(x : Even) ~(next x = x)
endspec
```

one possible model (out of many!) is the assignment of the even integers to Even, and of the function that increments an even number by 2 to next.

Each model has to *respect typing*; for example, given the above assignment to Even, the function that increments a number by 1 does not map all even numbers to even numbers, and therefore can not -- in the same model -- be assigned to next. Additionally, the claims (axioms, theorems and conjectures) of the spec have to be satisfied in the model. The axiom labeled nextEffect above states that the function assigned to op next maps any value of the type assigned to type-name Even to a

different value. So the squaring function, although type-respecting, could not be assigned to `next` since it maps 0 to itself.

If all type-respecting combinations of assignments of types to type-names and values to ops fail the one or more claims, the spec has no models and is called *inconsistent*. Although usually undesirable, an inconsistent spec is not by itself considered ill formed. The Specware system does not attempt to detect inconsistencies, but associated provers can sometimes be used to find them. Not always; in general it is undecidable whether a spec is consistent or not.

In general, the meaning of a construct in a model depends on the assignments of that model, and more generally on an *environment*: a model possibly extended with assignments to local-variables. For example, the meaning of the claim `fa(x : Even) ~(next x = x)` in axiom `nextEffect` depends on the meanings of `Even` and `next`, while the sub-expression `next x`, for example, also depends on an assignment (of an "even" value) to `x`. To avoid laborious phrasings, the semantic descriptions use language like "the function `next` applied to `x`" as shorthand for this lengthy phrase: "the function assigned in the environment to `next` applied to the value assigned in the environment to `x`".

When an environment is extended with an assignment to a local-variable, any assignments to synonymous ops or other local-variables are superseded by the new assignment in the new environment. In terms of Metaslang text, within the scope of the binding of local-variables, synonymous ops and earlier introduced local-variables (that is, having the same simple-name) are "hidden"; any use of the simple-name in that scope refers to the textually most recently introduced local-variable. For example, given:

```
def x = "op-x"
def y = let v = "let-v" in x
def z = let x = "let-x" in x
```

the value of `y` is `"op-x"` (op `x` is not hidden by the local-variable `v` of the let-binding), whereas the value of `z` is `"let-x"` (op `x` *is* hidden by the local-variable `x` of the let-binding).

2.1.3. Type-correctness

Next to the general requirement that each model respects typing, there are specific restrictions for various constructs that constrain the possible types for the components. For example, in an application `f(x)`, the type of the actual-parameter `x` has to match

the domain type of function `f`. These requirements are stated in the relevant sections of this language manual. If no type-respecting combinations of assignments meeting all these requirements exist for a given spec, it is considered *type-incorrect* and therefore *ill formed*. This is determined by Specware while elaborating the spec, and signaled as an error. Type-incorrectness differs from inconsistency in that the meaning of the claims does not come into play, and the question whether an ill-formed spec is consistent is moot.

To be precise, there are subtle and less subtle differences between type-incorrectness of a spec and its having no type-respecting combinations of assignments. For example, the following spec is type-correct but has no models:

```
spec
  type Empty = | Just Empty
  op IdoNotExist : Empty
endspec
```

The explanation is that the type-definition for `Empty` generates an *implicit* axiom that all inhabitants of the type `Empty` must satisfy, and for this recursive definition the axiom effectively states that such creatures can't exist: the type `Empty` is uninhabited. That by itself is no problem, but precludes a type-respecting assignment of an inhabitant of `Empty` to op `IdoNotExist`. So the spec, while type-correct, is actually inconsistent. See further under *Type-definitions*.

Here is a type-incorrect spec that has type-respecting combinations of assignments:

```
spec
  type Outcome = | Positive | Negative
  type Sign = | Positive | Zero | Negative
  def whatAmI = Positive
endspec
```

Here there are two constructors `Positive` of different types, the type `Outcome` and the type `Sign`. That by itself is fine, but when such "overloaded" constructors are used, the context must give sufficient information which is meant. Here, the use of `Positive` in the definition for op `whatAmI` leaves both possibilities open; as used it is *type-ambiguous*. Metaslang allows omitting type information provided that, given a type assignment to all local-type-variables in scope, unique types for all typed constructs, such as expressions and patterns, can be inferred from the context. If no complete and unambiguous type-assignment can be made, the spec is not accepted by the Specware system. Type-ambiguous expressions can be disambiguated by using a type annotation, as described under *Annotated-expressions*. In the example, the definition of `whatAmI` can be disambiguated in either of the following ways:

```
def whatAmI : Sign = Positive
def whatAmI = Positive : Sign
```

Also, if the spec elsewhere contains something along the lines of:

```
op signToNat : Sign -> Nat
def sw = signToNat whatAmI
```

that is sufficient to establish that `whatAmI` has type `Sign` and thereby disambiguate the use of `Positive`. See further under *Op-definitions* and *Structors*.

2.1.4. Constructive

When code is generated for a spec, complete "self-contained" code is only generated for type-definitions and op-definitions that are fully *constructive*. Non-constructiveness is "contagious": a definition is only constructive if all components of the definition are. The type of a type-name without definition is not constructive. A type is only constructive if all component types are. An op without definition is non-constructive, and so is an op whose type is non-constructive. A quantification is non-constructive. The polymorphic inbuilt-op = for testing equality and its inequality counterpart ~= are only constructive for *discrete types* (see below).

A type is called discrete if the equality predicate = for that type is constructive. The inbuilt and base-library types `Boolean`, `Integer`, `NonZeroInteger`, `Nat`, `PosNat`, `Char`, `String` and `Compare` are all discrete. Types `List` T and `Option` T are discrete when T is. All function types are non-discrete (even when the domain type is the unit type). Sum types, product types and record types are discrete whenever all component types are. Subtype `(`T `|` P`)` is discrete when supertype T is. (Predicate P need not be constructive: the equality test is that of the supertype.) Quotient type T / Q is discrete when predicate Q is constructive. (Type T need not be discrete: the equality test on the quotient type is just the predicate Q applied to pairs of members of the Q-equivalence classes.)

2.2. Lexical conventions

A Metaslang text consists of a sequence of symbols, possibly interspersed with whitespace. The term *whitespace* refers to any non-empty sequence of spaces, tabs, newlines, and comments (described below). A symbol is presented in the text as a sequence of one or more "marks" (ASCII characters). Within a composite (multi-mark) symbol, as well as within a unit-identifier, no whitespace is allowed, but whitespace may be needed between two symbols if the fisrt mark of the second symbol could be taken to be the continuation if the first symbol. More precisely, letting X, Y and Z stand for arbitrary (possibly empty) sequences of marks, and m for a single mark, then whitespace is required between two adjacent symbols, the first being X and the second mY, when for some Z the sequence XmZ is also a symbol. So, for example, whitespace is required where shown in `succ 0` and `op! :Nat->Nat`, since `succ0` and `!:` are valid symbols, but none is required in the formula `n+1`.

Inside literals (constant-denotations) whitespace is also disallowed, except for "significant-whitespace" as described under *String-literals*.

Other than that, whitespace -- or the lack of it -- has no significance. Whitespace can be used to lay-out the text for readability, but as far as only the meaning is concerned, the following two presentations of the same spec are entirely equivalent:

```
spec
  type Even
  op next : Even -> Even
  axiom nextEffect is
    fa(x : Even) ~(next x = x)
endspec

spec type   Even op   next : Even -> Even axiom nextEffect
is fa(x : Even)~(next     x              = x)endspec
```

2.2.1. Symbols and Names

symbol ::= simple-name | literal | special-symbol

simple-name ::= first-syllable { _ next-syllable } *

first-syllable ::= first-word-syllable | non-word-syllable

next-syllable `::=` next-word-syllable `|` non-word-syllable

first-word-syllable `::=` word-start-mark `{` word-continue-mark `}*`

next-word-syllable `::=` word-continue-mark `{` word-continue-mark `}*`

word-start-mark `::=` letter

word-continue-mark `::=` letter `|` decimal-digit `|` `'` `|` `?`

letter `::=`
```
    A | B | C | D | E | F | G | H | I | J | K | L | M
  | N | O | P | Q | R | S | T | U | V | W | X | Y | Z
  | a | b | c | d | e | f | g | h | i | j | k | l | m
  | n | o | p | q | r | s | t | u | v | w | x | y | z
```

decimal-digit `::=` `0` `|` `1` `|` `2` `|` `3` `|` `4` `|` `5` `|` `6` `|` `7` `|` `8` `|` `9`

non-word-syllable `::=` non-word-mark `{` non-word-mark `}*`

non-word-mark `::=`
```
    ` | ~ | ! | @ | $ | ^ | & | * | -
  | = | + | \ | | | : | < | > | / | ' | ?
```

special-symbol `::=` `_` `|` `(` `|` `)` `|` `[` `|` `]` `|` `{` `|` `}` `|` `;` `|` `,` `|` `.`

Sample simple-names:

```
    Date                    $$              ?!
    yymmdd2date             <*>             :=:
    well_ordered?           ~==             x'
    c_<+>                   /_47a
```

For convenience, here are the 13 printing ASCII marks that, next to letters and
decimal-digits, can *not* occur as a non-word-mark:

```
    #    %    "    _    ;    ,    .
    (    )    [    ]    {    }
```

Restriction. As mentioned before, no whitespace is allowed in symbols: while `anode`
is a single simple-name, both `a node` and `an ode` consist of two simple-names.
Further, the case (lower or upper) of letters in simple-names is significant:
`grandparent`, `grandParent` and `grandpaRent` are three different simple-names.

Restriction. In general, simple-names can be chosen freely. However, the following *reserved symbols* have a special meaning and must not be used for simple-names:

```
as          embed      from      let           quotient
axiom       embed?     generate  morphism      spec
by          endspec    if        obligations   the
case        ex         import    of            then
choose      ex1        in        op            theorem
conjecture  fa         infixl    project       true
def         false      infixr    prove         type
else        fn         is        qualifying    where


:           ||         |         =         <=>       <-        <<
::          &&         ~         ~=        =>        ->        +->
```

They each count as a single symbol, and no whitespace is allowed inside any reserved symbol. The symbols beginning with _ are for use with X-Symbol as described in the following section.

Non-word-syllables can be used to choose convenient simple-names for operators that, conventionally, are written with non-alphabetic marks.

Some Metaslang users follow the convention of using simple-names that start with a capital letter for unit-identifiers and type-names and for constructors, while simple-names chosen for ops and field-names start with a lower-case letter. Both plain local-variables and local-type-variables are often chosen to be single lower-case letters: x, y, z, a, b, c, with the start of the alphabet preferred for local-type-variables. Op-names of predicates (that is, having some type $T \rightarrow$ Boolean) often end with the mark ?. These are just conventions that users are free to follow or ignore, but in particular some convention distinguishing constructors from ops and local-variables is recommended.

2.2.2. X-Symbol

X-Symbol is an XEmacs package that allows ascii strings to encode non-ascii symbols. Specware files contain the ascii strings but when displayed in an XEmacs in with X-Symbol mode, the special symbol is shown. For example, the Specware symbol _forall is displayed as $\forall$ and _or is displayed as $\vee$. Many of the mathematical symbols have x-symbol aliases, for example, fa and _forall (displayed $\forall$ by X-Symbol) are alternatives for universal quantification, and _rightarrow (displayed $\longrightarrow$ or $\longrightarrow$ or by X-Symbol) may be used anywhere where -> is allowed. The user

may use non-reserved X-Symbol tokens as symbol names. The available X-Symbols are listed under the X-Symbol menu in XEmacs when X-Symbol mode is on.

2.2.3. Comments

comment ::= line-end-comment | block-comment

line-end-comment ::= **%** line-end-comment-body

line-end-comment-body ::=
     any-text-up-to-end-of-line

block-comment ::= **(*** block-comment-body ***)**

block-comment-body ::=
     any-text-including-newlines-and-nested-block-comments

Sample comments:

```
% keys must be unique
(* op yymmdd2Date : String -> Date *)
```

Metaslang allows two styles of comments. The %-style is light-weight, for adding comment on a line *after* the formal text (or taking a line on its own, but always confined to a single line). The (*...*)-style can be used for blocks of text, spanning several lines, or stay within a line. Any text remaining on the line after the closing *) is processed as formal text. Block-comments may be nested, so the pairs of brackets (* and *) must be balanced.

A block-comment can not contain a line-end-comment and vice versa: whichever starts first has "the right of way". For example, (* 100 % or more! *) is a block-comment with block-comment-body 100 % or more! . The % here is a mark like any other; it does not introduce a line-end-comment. Conversely, in the line-end-comment % op <*> stands for (*) the (* is part of the line-end-comment-body; it does not introduce a block-comment. Note also that % and (* have no special significance in literals (which must not contain whitespace, including comments): "100 % or more!" is a well-formed string-literal.

2.3. Units

A "unit" is an identifiable unit-term, where "identifiable" means that the unit-term can be referred to by a unit-identifier. Unit-terms can be "elaborated", resulting in specs, morphisms, diagrams or other entities. The effect of elaborating a unit-definition is that its unit-term is elaborated and becomes associated with its unit-identifier.

For the elaboration of a unit-term to be meaningful, it has to be well formed and result in a well-formed -- and therefore type-correct -- entity. Well-formedness is a stricter requirement than type-correctness. If a unit-term or one of its constituent parts does not meet any of the restrictions stated in this language manual, it is ill formed. This holds throughout, also if it is not mentioned explicitly for some syntactic construct. Well-formedness is more than a syntactic property; in general, to establish well-formedness it may be necessary to "discharge" (prove) proof obligations engendered by the unit-term.

A Specware project consists of a collection of Metaslang unit-definitions. They can be recorded in one or more Specware files. There are basically two styles for recording unit-definitions using Specware files. In the single-unit style, the file, when processed by Specware, contributes a single unit-definition to the project. In the multiple-unit style, the file may contribute several unit-definitions. The two styles may be freely mixed in a project (but not in the same Specware file). This is explained in more detail in what follows.

```
unit-definition  ::=  unit-identifier  =  unit-term

unit-term  ::=
        spec-term
      | morphism-term
      | diagram-term
      | target-code-term
      | proof-term

specware-file-contents  ::=
        unit-term
      | infile-unit-definition  {  infile-unit-definition  } *

infile-unit-definition  ::=  fragment-identifier  =  unit-term

fragment-identifier  ::=  simple-name
```

Unit-definitions may use other unit-definitions, including standard libraries, which in Specware 4.2 are supposed to be part of each project. However, the dependencies between units must not form a cycle; it must always be possible to arrange the unit-definitions in an order in which later unit-definitions only depend on earlier ones. How unit-definitions are processed by Specware is further dealt with in the Specware User Manual.

As mentioned above, unit-definitions are collected in Specware files, which in Specware 4.2 must have an `.sw` extension. The Specware files do not directly contain the unit-definitions that form the project. In fact, a user never writes unit-definition explicitly. These are instead determined from the contents of the Specware files using the following rules. There are two possibilities here. The first is that the specware-file-contents consists of a single unit-term. If `P.sw` is the path for the Specware file, the unit being defined has as its unit-identifier `P`. For example, if file `C:/units/Layout/Fixture.sw` contains a single unit-term `U`, the unit-identifier is `/units/Layout/Fixture`, and the unit-definition it contributes to the project is

```
  /units/Layout/Fixture = U
```

(Note that this is not allowed as an infile-unit-definition in a specware-file-contents, since the unit-identifier is not a fragment-identifier.)

The second possibility is that the Specware file contains one or more infile-unit-definitions. If `I` is the fragment-identifier of such an infile-unit-definition, and `P.sw` is the path for the Specware file, the unit being defined has as its unit-identifier `P#I`. For example, if file `C:/units/Layout/Cart.sw` contains an infile-unit-definition `Pos = U`, the unit-identifier is `/units/Layout/Cart#Pos`, and the unit-definition it contributes to the project is

```
  /units/Layout/Cart#Pos = U
```

2.3.1. Unit Identifiers

unit-identifier `::=` swpath-based-path | relative-path

swpath-based-path `::=` / relative-path

relative-path `::=` { path-element / }* path-element [# fragment-identifier]

path-element `::=` path-mark { path-mark }*

```
path-mark  ::=
       letter  |  decimal-digit
     |  !  |  $  |  &  |  ′  |  +  |  –
     |  =  |  @  |  ^  |  `  |  ~  |  .
```

Unit-identifiers are used to identify unit-terms, by identifying files in a file store that contain unit-terms or infile-unit-definitions. Which path-marks can actually be used in forming a path-element may depend on restrictions of the operating system. The path-elements .. and . have special meanings: "parent folder" and "this folder". Under than this use, the mark . should not be used as the first or last path-mark of a path-element.

Typically, only a final part of the full unit-identifier is used. When Specware is started with environment variable `SWPATH` set to a semicolon-separated list of pathnames for directories, the Specware files are searched for relative to these pathnames; for example, if `SWPATH` is set to `C:/units/Layout;C:/units/Layout/Cart`, then `C:/units/Layout/Fixture.sw` may be shortened to `/Fixture`, and `C:/units/Layout/Cart.sw` to `/Cart`. How unit-definitions are processed by Specware is further dealt with in the Specware User Manual.

Further, unit-identifiers can be relative to the directory containing the Specware file in which they occur. So, for example, both in file `C:/units/Layout/Fixture.sw` and in file `C:/units/Layout/Cart.sw`, unit-identifier `Tools/Pivot` refers to the unit-term contained in file `C:/units/Layout/Tools/Pivot.sw`, while `Props#SDF` refers to the unit-term of infile-unit-definition `SDF = ...` contained in file `C:/units/Layout/Props.sw`. As a special case, a unit-term with the same name as the file may be referenced without a fragment-identifier. For example, in the current case, if the file `C:/units/Layout/Props.sw` contains the unit-term of infile-unit-definition `Props = ...`, then this unit-term can be referred to either by `Props#Props` or `Props`.

The unit-identifier must identify a unit-definition as described above; the elaboration of the unit-identifier is then the result of elaborating the corresponding unit-term, yielding a spec, morphism, diagram, or other entity.

2.3.2. Specs

```
spec-term  ::=
       unit-identifier
     |  spec-form
```

> | spec-qualification
> | spec-translation
> | spec-substitution
> | diagram-colimit
> | obligator

Restriction. When used as a spec-term, the elaboration of a unit-identifier must yield a spec.

The elaboration of a spec-term, if defined, yields an "expanded" spec-form as defined in the next subsection.

2.3.2.1. Spec Forms

spec-form `::=` **spec** { declaration }* **endspec**

Sample spec-forms:

```
spec import Measures import Valuta endspec
```

An *expanded* spec-form is a spec-form containing no import-declarations.

The elaboration of a spec-form yields the Metaslang text which is that spec itself, after expanding any import-declarations. The *meaning* of that text is the class of models of the spec, as described throughout this Chapter.

2.3.2.2. Qualifications

Names of types and ops may be *simple* or *qualified*. The difference is that simple-names are "unqualified": they do not contain a dot sign ".", while qualified-names are prefixed with a "qualifier" followed by a dot. Examples of simple-names are `Date`, `today` and `<*>`. Examples of qualified-names are `Calendar.Date`, `Calendar.today` and `Monoid.<*>`.

Qualifiers can be used to disambiguate. For example, there may be reason to use two different ops called `union` in the same context: one for set union, and one for bag (multiset) union. They could then more fully be called `Set.union` and `Bag.union`, respectively. Unlike in earlier versions of Specware, there is no rigid relationship between qualifiers and the unit-identifiers identifying specs. The author of a

collection of **specs** may use the **qualifier** deemed most appropriate for any **type-** or **op-name**. For example, there could be a single **spec** dubbed `SetsAndBags` that introduces two new **ops**, one called `Set.union` and one called `Bag.union`. Generally, types and **ops** that "belong together" should receive the same **qualifier**. It is up to the author of the **specs** to determine what belongs together.

Type-names and **ops** are *introduced* in a **declaration** or **definition**, and may then be *employed* elsewhere in the same **spec**. Thus, all occurrences of a **type-** or **op-name** can be divided into "introductions" and "employs". The **name** as introduced in an introduction is the *full* **name** of the type or **op**. If that **name** is a **simple-name**, the full **name** is a **simple-name**. If the **name** as introduced is a **qualified-name**, then so is the full **name**.

For employs the rules are slightly different. First, if the **name** employed occurs just like that in an introduction, then it is the full **name**. Also, if the **name** employed is qualified, it is the full **name**. Otherwise, the **name** as employed may be unqualified shorthand for a qualified full **name**. For example, given an employ of the unqualified **type-name** `Date`, possible qualified full **names** for it are `Calendar.Date`, `DateAndTime.Date`, `Diary.Date`, and so on. But, of course, the full **name** must be one that is introduced in the **spec**. If there is precisely one **qualified-name** introduced whose last part is the same as the **simple-name** employed, then that **name** is the full **name**. Otherwise, type information may be employed to disambiguate ("resolve overloading").

Here is an illustration of the various possibilities:

```
spec
  type Apple
  type Fruit.Apple
  type Fruit.Pear
  type Fruit.Date
  type Calendar.Date
  type Fruit.Basket = Apple * Pear * Date
endspec
```

In the definition of type `Fruit.Basket` we have three unqualified employs of **type-names**, viz. `Apple`, `Pear` and `Date`. The **name** `Apple` is introduced like that, so the employ `Apple` already uses the full **name**; it does not refer to `Fruit.Apple`. The **name** `Pear` is nowhere introduced just like that, so the employ must be shorthand for some qualified full **name**. There is only one applicable introduction, namely `Fruit.Pear`. Finally, for `Date` there are two candidates: `Fruit.Date` and `Calendar.Date`. This is ambiguous, and in fact an error. To correct the error, the

employ of `Date` should be changed into either `Fruit.Date` or `Calendar.Date`, depending on the intention.

It is possible to give a qualification in one go to all **simple-names** introduced in a **spec**. If `Q` is a **qualifier**, and `S` is a term denoting a **spec**, then the term `Q` `qualifying` `S` denotes the same **spec** as `S`, except that each introduction of an **simple-name** `N` is replaced by an introduction of the **qualified-name** `Q.N`. Employs that before referred to the unqualified introduction are also accordingly qualified, so that they now refer to the qualified introduction. For example, the value of

```
Company qualifying spec
  type Apple
  type Fruit.Apple
  type Fruit.Pear
  type Fruit.Basket = Apple * Pear
endspec
```

is the same as that of

```
spec
  type Company.Apple
  type Fruit.Apple
  type Fruit.Pear
  type Fruit.Basket = Company.Apple * Fruit.Pear
endspec
```

spec-qualification ::= qualifier **qualifying** spec-term

qualifier ::= simple-name

name ::= simple-name | qualified-name

qualified-name ::= qualifier **.** simple-name

Restriction. **qualifying** is a special kind of translation. As such, it is not allowed to rename a **type** to have the same name as a previously existing (and different) **type**, or to rename an **op** to have the same name as a previously existing (and different) **op**.

Sample **names**:

```
  Key              $
  Calendar.Date    Monoid.<*>
```

Sample spec-qualification:

```
Weight qualifying /Units#Weights
```

Let R be the result of elaborating spec-term S. Then the elaboration of qualification Q qualifying S, where Q is a qualifier, is R with each unqualified type-name, op-name or claim-name N introduced there replaced by the qualified-name $Q.N$. The same replacement applies to all employs of N identifying that introduced simple-name. As always, the result of the replacement is required to be a well-formed spec.

For example, the elaboration of

```
Buffer qualifying spec
   op size : Nat
   axiom LargeSize is size >= 1024
endspec
```

results in:

```
spec
   op Buffer.size : Nat
   axiom Buffer.LargeSize is Buffer.size >= 1024
endspec
```

Because of the restriction on collisions, the following is illegal, as f in the original spec would be renamed to collide with the pre-existing (and distinct!) X.f.

```
X qualifying spec
   op  X.f : Nat
   def f = 3
endspec
```

However, it is legal to qualify one alias for a type or op in such a way as to collapse it onto another alias for the same type or op, as in the following

```
X qualifying spec
   op  {f, X.f} : Nat
endspec
```

which results in

```
spec  op X.f : Nat endspec
```

2.3.2.3. Translations

spec-translation ::= **translate** spec-term **by** name-map

name-map ::= **{** [name-map-item **{ ,** name-map-item **}***] **}**

name-map-item ::=
      type-name-map-item
   | op-name-map-item
   | wildcard-map-item

type-name-map-item ::= [**type**] name **+->** name

op-name-map-item ::= [**op**] annotable-name **+->** annotable-name

annotable-name ::= name [**:** type-descriptor]

wildcard-map-item ::= wildcard **+->** wildcard

wildcard ::= simple-wildcard | qualified-wildcard

simple-wildcard ::= _

qualified-wildcard ::= qualifier **.** simple-wildcard

Restriction. The name-map of a spec-translation may not contain more than one name-map-item pertaining to the same type-name or the same op-name in the spec resulting from elaborating the spec-term. For example, the following is not a lawful spec-translation:

```
translate spec type T by {T +-> A, T +-> B}
```

Restriction. A spec-translation may not map two different type-names or two different op-names to the same simple-name. Note that this implies that type- and op-names cannot be translated to simple-names defined in the base libraries.

Sample spec-translation:

```
translate A by {Counter +-> Register, tally +-> incr}
```

Let R be the result of elaborating spec-term S. In elaborating a spec-translation, first any wildcard-map-items are expanded as follows. A simple-wildcard matches each simple-name that is a type-name or op-name of S. A qualified-wildcard $Q._$ matches each qualified-name having the same qualifier Q that is a type-name or

op-name of S. A wildcard-map-item $W0$ `+->` $W1$ is expanded into a list of name-map-items not containing wildcards by taking each name N matched by $W0$, and replacing both simple-wildcards occurring in $W0$ `+->` $W1$ by the simple-name of N, that is, N with a possible qualification stripped off. After expansion, the elaboration of `translate` S by $\{\ M_1\ $ `+->` $\ N_1,\ \ldots\ M_n\ $ `+->` $\ N_n\ \}$ is R with each occurrence of a name M_i replaced by N_i. All other names are mapped to themselves, i.e., they are unchanged. The presence of a type annotation in a name-map-item, as in `X:E +->` `cross`, indicates that the name-map-item refers to an op-name; additionally, on the left-hand side such an annotation may serve to disambiguate between several synonymous ops, and then there must be an op in R of the type given by the type-descriptor. If the right-hand side of a name-map-item carries a type annotation, its type-descriptor must conform to the type of the op-name in the resulting spec. Without such annotation on either side, if a name to be translated is introduced both as a type-name and as an op-name in R, it must be preceded by `type` or `op` to indicate which of the two is meant. Otherwise the indication `type` or `op` is not required, but allowed; if present, it must correspond to the kind of simple-name (type-name or op-name) to be translated.

For example, the elaboration of

```
translate spec
  type E
  op i : E
endspec by {
  E +-> Counter,
  i +-> reset
}
```

results in:

```
spec
  type Counter
  op reset : Counter
endspec
```

To illustrate the use of wildcards: The elaboration of

```
translate spec
  type M.Length
  op M.+ infixl 25 : M.Length * M.Length -> M.Length
endspec by {M._ +-> Measure._}
```

results in this spec:

```
spec
  type Measure.Length
  op Measure.+ infixl 25 :
             Measure.Length * Measure.Length -> Measure.Length
endspec
```

A spec-qualification `Q qualifying S` is convenient shorthand for the spec-translation `translate S by {_ +-> Q._}`.

2.3.2.4. Substitutions

spec-substitution `::=` spec-term `[` morphism-term `]`

Sample spec-substitution:

```
Routing#Basis[morphism /Coll/Lattice ->
                     /Coll/LatticeWithTop {} ]
```

The elaboration of spec-substitution `S[M]` yields the spec `T` obtained as follows. Let spec `R` be the result of elaborating `S`, and morphism `N` that of `M`. Let specs `D` and `C` be the domain and codomain of `N`. First, remove from `R` all declarations of `D`, and subject the result to the effect of `N`, meaning that all name translations of `N` and all extensions with declarations are performed. Then add the declarations of `C`, but without duplications, i.e., as if `C` is imported. The result obtained is `T`.

Restriction. Spec `D` must be a "sub-spec" of spec `R`, meaning that each declaration of `D` is also a declaration of `R`.

Informally, `T` is to `R` as `C` is to `D`.

Except when `R` introduces, next to the type- and op-names it has in common with `D`, new type- or op-names that also occur in `C`, the result spec `T` is a categorical colimit of this pushout diagram:

```
D ---------> R
|            |
|            |
|            |
v            v
C ---------> T
```

Although isomorphic to the result that would be obtained by using a diagram-colimit, T is more "user-oriented" in two ways: the names in T are names from C, and claims of D not repeated in C are not repeated here either.

For example, assume we have:

```
A = spec
  type Counter
  op reset: Counter
  op tally : Counter -> Counter
  axiom Effect is
    fa (c : Counter) ~(tally c = c)
endspec

B = spec
  type Register = Nat
  def reset = 0
  def incr c = c+1
endspec

M = morphism A -> B {Counter +-> Register, tally +-> incr}

AA = spec
  import A
  type Interval = {start: Counter, stop: Counter}
  op isEmptyInterval? : Interval -> Boolean
  def isEmptyInterval? {start = x, stop = y} = (x = y)
endspec
```

Then the result of `AA[M]` is the same as that of this spec:

```
spec
  import B
  type Interval = {start: Register, stop: Register}
  op isEmptyInterval? : Interval -> Boolean
  def isEmptyInterval? {start = x, stop = y} = (x = y)
endspec
```

2.3.2.5. Diagram Colimits

diagram-colimit ::= **colimit** diagram-term

The result of elaborating a diagram-colimit is the spec which is the apex of the cocone forming the colimit in the category of specs and spec-morphisms. As always, the result is required to be well formed. See further the Specware Tutorial.

2.3.2.6. Obligators

obligator ::= **obligations** unit-term

Restriction. The unit-term of an obligator must either be a spec-term or a morphism-term.

The result of elaborating an obligator is a spec containing the proof obligations engendered by the spec or morphism resulting from elaborating its unit-term. These proof obligations are expressed as conjectures; they can be discharged by proving them, using proof-terms. See further the Specware User Manual.

2.3.3. Morphisms

morphism-term ::=
        unit-identifier
    | spec-morphism

spec-morphism ::= **morphism** spec-term **->** spec-term name-map

A morphism is a formal mapping between two expanded spec-forms that describes exactly how one is translated or extended into the other. It consists of the two specs, referred to as "domain" and "codomain", and a mapping of all type- and op-names introduced in the domain spec to type- and op-names of the codomain spec. To be well-formed, a morphism must obey conditions that express that it is a proper refinement of the domain into the codomain ' spec.

Restriction. When used as a morphism-term, the elaboration of a unit-identifier must yield a morphism.

Restriction ("proof obligations"). Given spec-morphism morphism S -> T { M_1 +-> N_1, ... M_n +-> N_n } let R be the result of elaborating S, and let S' be R with each occurrence of a name M_i replaced by N_i. The same rules apply as for spec-translation translate S by {...}, and the result S' must be well formed, with the exception that the restriction on spec-translations requiring different type- or

op-names to be mapped to different simple-names does not apply here. Let T' be the result of elaborating T. Then, first, each type-name or op-name introduced in S' must also be introduced in T'. Further, no type-name or op-name originating from a library spec may have been subject to translation. Finally, each claim in S' must be a theorem that follows from the claims of T'. Collectively, the claims in S' are known as the *proof obligations* engendered by the morphism. They are the formal expression of the requirement that the step from S' to T' is a proper refinement.

For example, in

```
S = spec endspec
T = spec type Bullion = (Char | isAlpha) endspec
M = morphism S -> T {Boolean -> Bullion}
```

the type-name `Boolean`, which originates from a library spec, is subject to translation. Therefore, `M` is not a proper morphism. Further, in

```
S = spec
        op f : Nat -> Nat
        axiom ff is fa(n:Nat) f(n) ~= f(n+1)
    endspec

T = spec
        def f(n:Nat) = n*n rem 5
    endspec

M = morphism S -> T
```

axiom `ff` does not follow from (the axiom implied by) the op-definition for `f` in spec `T`, since `f(2) = f(3) = 4`. Therefore, `M` is not a proper morphism here either.

Sample spec-morphism:

```
morphism A -> B {Counter +-> Register, tally +-> incr}
```

The elaboration of spec-morphism `morphism` S `->` T `{`M`}` results in the morphism whose domain and codomain are the result of elaborating S and T, respectively, and whose mapping is given by the list of name-map-items in M, using type annotations and indicators `type` and `op` as described for spec-translations, and extended to all domain type- and op-names not yet covered by having these map to themselves. (In particular, simple-names from the base-libraries always map to themselves.)

2.3.4. Diagrams

```
diagram-term  ::=
        unit-identifier
      | diagram-form
```

```
diagram-form  ::= diagram { diagram-element { , diagram-element }* }
```

```
diagram-element  ::=
        diagram-node
      | diagram-edge
```

```
diagram-node  ::= simple-name +-> spec-term
```

```
diagram-edge  ::= simple-name : simple-name -> simple-name +-> morphism-term
```

Restriction. When used as a diagram-term, the elaboration of a unit-identifier must yield a diagram.

Restriction. In a diagram, the first simple-name of each diagram-node and diagram-edge must be unique (i.e., not be used more than once in that diagram). Further, for each diagram-edge E : ND -> NC +-> M, there must be diagram-nodes ND +-> D and NC +-> C of the diagram such that, after elaboration, M is a morphism from D to C.

Sample diagram:

```
diagram {
    A           +-> /Coll/Lattice,
    B           +-> /Coll/LatticeWithTop,
    m : A -> B +-> /Coll/AddTop,
    C           +-> Routing#Basis,
    i : A -> C +-> morphism /Coll/Lattice ->
                                Routing#Basis {}
}
```

The result of elaborating a diagram-form is the categorical diagram whose nodes are labeled with the specs and whose edges are labeled with the morphisms that result from elaborating the corresponding spec-terms and morphism-terms.

2.3.5. Target Code Terms

```
target-code-term ::=
        generate target-language-name spec-term [ generate-option ]

generate-option ::=
        in string-literal | with unit-identifier

target-language-name ::= c | java | lisp
```

Sample target-code-term:

```
generate lisp /Vessel#Contour
                    in "C:/Projects/Vessel/Contour.lisp"
```

The elaboration of a target-code-term for a well-formed spec-term generates code in the language suggested by the target-language-name (currently only C, Java, and Common Lisp); see further the Specware User Manual.

2.3.6. Proof Terms

```
proof-term ::=
        prove claim-name in spec-term
                            [ with prover-name ]
                            [ using claim-list ]
                            [ options prover-options ]

prover-name ::= snark

claim-list ::= claim-name { , claim-name }*

prover-options ::= string-literal
```

Restriction. The claim-names must occur as claim-names in the spec that results from elaborating the spec-term.

Sample proof-term:

```
prove Effect in obligations M
                        options "(use-paramodulation t)"
```

The elaboration of a proof-term invokes the prover suggested by the prover-name (currently only SNARK). The property to be proved is the claim of the first claim-name; the claim-list lists the hypotheses (assumptions) that may be used in the proof. The prover-options are prover-specific and are not further described here. For details, see the Specware User Manual.

2.4. Declarations

```
declaration ::=
        import-declaration
      | type-declaration
      | op-declaration
      | definition

definition ::=
        type-definition
      | op-definition
      | claim-definition

equals ::= is | =
```

Sample declarations:

```
import Lookup
type Key
op present : Database * Key -> Boolean
type Key = String
def present(db, k) = embed? Some (lookup (db, k))
axiom norm_idempotent is fa(x) norm (norm x) = norm x
```

2.4.1. Import-declarations

A spec may contain one or more import-declarations. On elaboration, these are "expanded". The effect is as if the bodies of these imported specs (themselves in elaborated form, which means that all import-declarations have been expanded, all

translations performed and all shorthand employs of names have been resolved to full names, after which only declarations or definitions of types, ops and claims are left) is inserted in place in the receiving spec.

For example, the result of

```
spec
  import spec
            type A.Z
            op b : Nat -> Z
         end
  type A.Z = String
  def b = toString
endspec
```

is this "expanded" spec:

```
spec
  type A.Z
  op b : Nat -> A.Z
  type A.Z = String
  def b = toString
endspec
```

For this to be well formed, the imported specs must be well formed by themselves; in addition, the result of expanding them in place must result in a well-formed spec.

There are a few restrictions, which are meant to catch unintentional naming mistakes. First, if two different imported specs each introduce a type or op with the same (full) name, the introductions must be identical declarations or definitions, or one may be a declaration and the other a "compatible" definition. For example, given

```
S1 = spec op e : Integer end
S2 = spec op e : Char end
S3 = spec def e = 0 end
```

the specs S1 and S3 can be imported together, but all other combinations of two or more co-imported specs result in an ill-formed spec. This restriction is in fact a special case of the general requirement that import expansion must result in a well-formed spec. Secondly, a type-name introduced in any of the imported specs cannot be re-introduced in the receiving spec except for the case of an "imported" declaration together with a definition in the receiving spec. Similarly for op-names, with the addition that an op-definition in the receiving spec must be compatible with

an op-declaration for the same name in an imported spec. The latter is again a special case of the general requirement that import expansion must result in a well-formed spec.

What is specifically excluded by the above, is giving a definition of a type or op in some spec, import it, and then redefining or declaring that type or op with the same full name in the receiving spec.

import-declaration ::= **import** spec-term

Sample import-declarations

```
import Lookup
```

An import-declaration is contained in some spec-form, and to elaborate that spec-form the spec-term of the import-declaration is elaborated first, giving some spec *S*. The import-declaration has then the effect as if the declarations of the imported spec *S* are expanded in place. This cascades: if spec *A* imports *B*, and spec *B* imports *C*, then effectively spec *A* also imports *C*. An important difference with earlier versions of Specware than version 4 is that multiple imports of the same spec have the same effect as a single import.

If spec *A* is imported by *B*, each model of *B* is necessarily a model of *A* (after "forgetting" any simple-names newly introduced by *B*). So *A* is then refined by *B*, and the morphism from *A* to *B* is known as the "import morphism". As it does not involve translation of type- or op-names, it can be denoted by morphism *A* -> *B* {}.

2.4.2. Type-declarations

type-declaration ::= **type** type-aliases [formal-type-parameters]

type-aliases ::= type-name | **{** type-name **,** type-name { **,** type-name }* **}**

formal-type-parameters ::= local-type-variable | **(** local-type-variable-list **)**

local-type-variable ::= simple-name

local-type-variable-list ::= local-type-variable { **,** local-type-variable }*

Restriction. Each local-type-variable of the formal-type-parameters must be a different simple-name.

Sample type-declarations:

```
type Date
type Array a
type Map(a, b)
type {Unit,Singleton} (a)
```

Every type-name used in a spec must be declared (in the same spec or in an imported spec, included the "base-library" specs that are always implicitly imported). A type-name may have *type parameters*. Given the example type-declarations above, some valid type-descriptors that can be used in this context are `Array Date`, `Array (Array Date)` and `Map (Nat, Boolean)`.

In a model of the spec, a type is assigned to each unparameterized type-name, while an infinite *family* of types is assigned to parameterized type-names "indexed" by tuples of types, that is, there is one family member, a type, for each possible assignment of types to the local-type-variables. So for the above example type-declaration of `Array` one type must be assigned to `Array Nat`, one to `Array Boolean`, one to `Array (Array Date)`, and so on. These assigned types could all be the same type, or perhaps all different, as long as the model respects typing.

The ability to have aliases ia provided primarily to support the colimit operation, in which case multiple differently named types may be glued together, with no obvious rule for giving one name precedence over any other. They may also be used to define ASCII and X-Symbol synonyms, etc.

2.4.3. Type-definitions

```
type-definition ::=
      type-abbreviation
    | new-type-definition

type-abbreviation    ::= type type-aliases [ formal-type-parameters ] equals
                              type-descriptor

new-type-definition ::= type type-aliases [ formal-type-parameters ] equals
                              new-type-descriptor
```

Sample type-abbreviations:

```
type Date = {year : Nat, month : Nat, day : Nat}
type Array a = List a
type Map(a, b) = (Array (a * b) | key_uniq?)
```

Sample new-type-definitions:

```
type Tree         a = | Leaf a | Fork (Tree a * Tree a)
type {Bush,Shrub} a = | Leaf a | Fork (Tree a * Tree a)
type Z3 =  Nat / (fn (m, n) -> m rem 3 = n rem 3)
```

In each model, the type assigned to the type-name of a type-abbreviation must be the same as the right-hand-side type-descriptor, while that assigned to the type-name of a new-type-definition must be isomorphic to the type of the new-type-descriptor. Thus while `Tree Nat` and `Bush Nat` from the examples are equivalent, they are not necessarily equal, but `Bush Nat` and `Shrub Nat` are truly synonymous.

For parameterized types, this extends to all possible assignments of types to the local-type-variables, taking the right-hand type-descriptors and new-type-descriptors as interpreted under each of these assignments. So, for the example, `Map(Nat, Char)` is the same type as `(Array (Nat * Char) | key_uniq?)`, and so on.

With *recursive* type-definitions, there are additional requirements. For example, consider

```
type Stack a =
  | Empty
  | Push {top : a, pop : Stack a}
```

This means that for each type `a` there is a value `Empty` of type `Stack a`, and further a function `Push` that maps values of type `{top : a, pop : Stack a}` to `Stack a`. Furthermore, the type assigned to `Stack a` must be such that all its inhabitants can be constructed *exclusively* and *uniquely* in this way: there is one inhabitant `Empty`, and all others are the result of a `Push`. Finally -- this is the point -- the type in the model must be such that its inhabitants can be constructed this way in *a finite number of steps*. So there can be no "bottom-less" stacks: deconstructing a stack using

```
def [a] hasBottom? (s : Stack a) : Boolean =
  case s of
      | Empty -> true
      | Push {top, pop = rest} -> hasBottom? rest
```

is a procedure that is guaranteed to terminate, always resulting in `true`.

In general, type-definitions generate implicit axioms, which for recursive definitions imply that the type is not "larger" than necessary. In technical terms, in each model the type is the least fixpoint of a recursive domain equation.

2.4.4. Op-declarations

op-declaration ::=

      op [type-variable-binder] formal-decl-expression [fixity] **:** type
       [equals expression]

    | **op** formal-decl-expression [fixity] **:** type-variable-binder type
       [equals expression]

type-variable-binder ::= `[ local-type-variable-list `]

formal-decl-expression ::= op-aliases | formal-decl-application

formal-decl-application ::= formal-decl-application-head formal-parameter

formal-decl-application-head ::= op-aliases | formal-decl-application

op-aliases ::= op-name | { op-name , op-name { , op-name }* }

formal-parameter ::= closed-pattern

fixity ::= associativity priority

associativity ::= **infixl** | **infixr**

priority ::= nat-literal

Note that the formal-decl-application of an op-declaration always uses prefix notation, even for infix-operators.

Sample op-declarations:

```
op usage : String

op [a,b,c] o infixl 24 :         (b -> c) * (a -> b) -> a -> c
op          o infixl 24 : [a,b,c] (b -> c) * (a -> b) -> a -> c
```

```
op f (s : String) ((i, j) : Nat * Nat) (c : Char) : String =
  substring (s, i, j) ^ toString c

op {Moon,Luna} : Satellite
```

An op-declaration introduces an op with an associated type. The type can be "monomorphic", like `String`, or "polymorphic" (indicated by a type-variable-binder). In the latter case, an indexed family of values is assigned to parameterized type-names "indexed" by tuples of types, that is, there is one family member, a typed value, for each possible assignment of types to the local-type-variables of the type-variable-binder, and the type of that value is the result of the corresponding substitution of types for local-type-variables on the polymorphic type of the op. In the examples above, the two forms given for the declaration of polymorphic `o` are entirely equivalent; they can be thought of as introducing a family of (fictitious) ops, one for each possible assignment to the local-type-variables `a`, `b` and `c`:

$o_{\text{Nat,String,Char}}$ `: (String -> Char) * (Nat -> String) -> Nat -> Char`

$o_{\text{Nat,Nat,Boolean}}$ `: (Nat -> Boolean) * (Nat -> Nat) -> Nat -> Boolean`

$o_{\text{Char,Boolean,Nat}}$ `: (Boolean -> Nat) * (Char -> Boolean) -> Char -> Nat`

and so on. Any op-definition for `o` must be likewise accommodating.

Only binary ops (those having some type `S * T -> U`) may be declared with a fixity. When declared with a fixity, the op may be used in infix notation, and then it is called an infix-operator. For `o` above, this means that `o(f, g)` and `f o g` may be used, interchangeably, with no difference in meaning. If the associativity is `infixl`, the infix-operator is called *left-associative*; otherwise, if the associativity is `infixr`, it is called *right-associative*. If the priority is `priority N`, the operator is said to have *priority N*. The nat-literal `N` stands for a natural number; if infix-operator `O1` has priority `N1`, and `O2` has priority `N2`, with `N1 < N2`, we say that `O1` has *lower priority* than `O2`, and that `O2` has *higher priority* than (or *takes priority over*) `O1`. For the role of the associativity and priority, see further at *Infix-applications*.

2.4.5. Op-definitions

```
op-definition ::=
        def [ type-variable-binder ] formal-def-expression [ : type-descriptor ] equals
            expression
     |  def formal-def-expression [ : [ type-variable-binder ] type-descriptor ] equals
            expression

formal-def-expression ::= op-name | formal-def-application

formal-def-application ::= formal-def-application-head formal-parameter

formal-def-application-head ::= op-name | formal-def-application
```

Sample op-definitions:

```
def usage = "Usage: Lookup key [database]"

def [a,b,c] o(f : b -> c, g: a -> b) : a -> c =
  fn (x : a) -> f(g x)

def o : [a,b,c] (b -> c) * (a -> b) -> a -> c =
  fn (f, g) -> fn (x) -> f(g x)

def o(f, g) x = f(g x)
```

Restriction. See the restriction under *Op-declarations* on redeclaring/redefining ops.

Note that a formal-def-expression always contains precisely one op-name, which is the op *being defined* by the op-definition. Note further that the formal-def-application of an op-definition always uses prefix notation, even for infix-operators.

In Specware 4.2, an op still may be defined without having been declared, but *this is now a deprecated feature*. When an op is defined without having been declared, the op-definition generates an implicit op-declaration for the op, provided a monomorphic type for the op can be unambiguously determined from the op-definition together with the uses of the op in applications and other contexts. In general, typing information in an op-definition may be omitted, but sufficient information must be supplied when used, so that all expressions can be assigned a type in the context in which they occur while uniquely associating the op with an op-declaration or op-definition. If two different associations both give type-correct specs, the spec is ambiguous and ill formed.

This feature may not persist in future releases, so users are advised to provide an op-declaration somewhere before each op-definition, either in the same spec or (more typically) in an imported spec. Alternatively, the newly expanded syntax for an *op-declaration* makes it simple to both declare and define an op within one declaration.

As for op-definitions, the presence of a type-variable-binder signals that the op being defined is polymorphic. In a model of the spec, an indexed family of typed values is assigned to a polymorphic op, with one family member for each possible assignment of types to the local-type-variables of the type-variable-binder, and the type of that value is the result of the corresponding type-instantiation for the polymorphic type of the op. Thus, we can reduce the meaning of a polymorphic op-definition to a family of (fictitious) monomorphic op-definitions.

An op-definition with formal-prefix-application

```
def H P = E
```

in which H is a formal-application-head, P is a formal-parameter and E an expression, is equivalent to the op-definition

```
def H = fn P -> E
```

For example,

```
def o (f, g) x = f(g x)
```

is equivalent to

```
def o (f, g) = fn x -> f(g x)
```

which in turn is equivalent to

```
def o = fn (f, g) -> fn x -> f(g x)
```

By this deparameterizing transformation for each formal-parameter, an equivalent unparameterized op-definition is reached. The semantics is described in terms of such op-definitions.

In each model, the typed value assigned to the op being defined must be the same as the value of the right-hand-side expression. For polymorphic op-definitions, this extends to all possible assignments of types to the local-type-variables.

An op-definition can be thought of as a special notation for an axiom. For example,

```
def [a] double (x : a) = (x, x)
```

can be thought of as standing for:

```
op double : [a] a -> a * a

axiom double_def is
   [a] fa(x : a) double x = (x, x)
```

In fact, Specware generates such axioms for use by provers. But in the case of recursive definitions, this form of axiomatization does not adequately capture the meaning. For example,

```
def f (n : Nat) : Nat = 0 * f n
```

is an improper definition, while

```
axiom f_def is
    fa(n : Nat) f n = 0 * f n
```

characterizes the function that maps every natural number to 0. The issue is the following. Values in models can not be *undefined* and functions assigned to ops must be *total*. But in assigning a meaning to a recursive op-definition, we -- temporarily -- allow *undefined* and partial functions (functions that are not everywhere defined on their domain type) to be assigned to recursively defined ops. In the thus extended class of models, the recursive ops must be the least-defined solution to the "axiomatic" equation (the least fixpoint as in domain theory), given the assignment to the other ops. For the example of f above this results in the everywhere undefined function, since 0 times *undefined* is *undefined*. If the solution results in an undefined value or a function that is not total (or for higher-order functions, functions that may return non-total functions, and so on), the op-definition is improper. Although Specware 4.2 does attempt to generate proof obligations for this condition, it currently covers only "simple" recursion, and not mutual recursion or recursion introduced by means of higher-order functions.

Functions that are determined to be the value of an expression, but that are not assigned to ops, need not be total, but the context must enforce that the function can not be applied to values for which it is undefined. Otherwise, the spec is ill formed.

2.4.6. Claim-definitions

claim-definition ::= claim-kind claim-name **is** claim

claim-kind ::= **axiom** | **theorem** | **conjecture**

claim-name ::= name

claim ::= [type-variable-binder] expression

Sample claim-definitions:

```
axiom norm_idempotent is
  norm o norm = norm

theorem o_assoc is
  [a,b,c,d] fa(f : c -> d, g : b -> c, h : a -> b)
    f o (g o h) = (f o g) o h

conjecture pivot_hold is
  let p = pivot hold in
    fa (n : {n : Nat | n < p}) ~(hold n = hold p)
```

Restriction. The type of the claim must be `Boolean`.

Restriction. A claim must not be an expression whose first symbol is `[`. In order to use such an expression as a claim, it can be parenthesized, as in

```
axiom nil_fits_nil is ([] fits [])
```

This restriction prevents ambiguities between claims with and without type-variable-binders.

When a type-variable-binder is present, the claim is polymorphic. A polymorphic claim may be thought of as standing for an infinite family of monomorphic claims, one for each possible assignment of types to the local-type-variables.

The claim-kind `theorem` should only be used for claims that have actually been proved to follow from the (explicit or implicit) axioms. In other words, giving them axiom status should not change the class of models. Theorems can be used by provers.

Conjectures are meant to represent proof obligations that should eventually attain theoremhood. Like theorems, they can be used by provers. This is only sound if

circularities (vicious circles) are avoided. This kind of claim is usually created automatically by the elaboration of an obligator, but can also be created manually.

The Specware system passes on the claim-name of the claim-definition with the claim for purposes of identification. Both may be transformed to fit the requirements of the prover, and appear differently there. Not all claims can be faithfully represented in all provers, and even when they can, the logic of the prover may not be up to dealing with them.

Remark. It is a common mistake to omit the part "claim-name is" from a claim-definition. A defensive style against this mistake is to have the claim always start on a new text line. This is additionally recommended because it may become required in future revisions of Metaslang.

2.5. Type-descriptors

```
type-descriptor  ::=
        type-arrow
     |  slack-type-descriptor

new-type-descriptor  ::=
        type-sum
     |  type-quotient

slack-type-descriptor  ::=
        type-product
     |  tight-type-descriptor

tight-type-descriptor  ::=
        type-instantiation
     |  closed-type-descriptor

closed-type-descriptor  ::=
        type-name
     |  Boolean
     |  local-type-variable
     |  type-record
     |  type-restriction
     |  type-comprehension
     |  ( type-descriptor )
```

(The distinctions "**slack-**", "**tight-**" and "**closed-**" before "**type-descriptor**" have no semantic significance. The distinction merely serves the purpose of diminishing the need for parenthesizing in order to avoid grammatical ambiguities.)

Sample **type-descriptors**:

```
List String * Nat -> Option String
a * Order a * a
PartialFunction (Key, Value)
Key
Boolean
a
{center : XYpos, radius : Length}
(Nat | even)
{k : Key | present (db, k)}
(Nat * Nat)
```

Sample **new-type-descriptors**:

```
| Point XYpos | Line XYpos * XYpos
Nat / (fn (m, n) -> m rem 3 = n rem 3)
```

The meaning of the **type-descriptor** `Boolean` is the "inbuilt" type inhabited by the two logical (truth) values `true` and `false`. The meaning of a parenthesized **type-descriptor** `(T)` is the same as that of the enclosed **type-descriptor** `T`.

The various other kinds of **type-descriptors** and **new-type-descriptors** not defined here are described each in their following respective sections, with the exception of **local-type-variable**, whose (lack of) meaning as a **type-descriptor** is described below.

Restriction. A **local-type-variable** may only be used as a **type-descriptor** if it occurs in the scope of a **formal-type-parameters** or **type-variable-binder** in which it is introduced.

Disambiguation. A single **simple-name** used as a **type-descriptor** is a **local-type-variable** when it occurs in the scope of a **formal-type-parameters** or **type-variable-binder** in which it is introduced, and then it identifies the textually most recent introduction. Otherwise, the **simple-name** is a **type-name**.

A **local-type-variable** used as a **type-descriptor** has no meaning by itself, and where relevant to the semantics is either "indexed away" (for parameterized types) or "instantiated away" (when introduced in a **formal-type-parameters** or **type-variable-binder**) before a meaning is ascribed to the construct in which it occurs. Textually, it has a scope just like a plain **local-variable**.

2.5.1. Type-sums

    type-sum ::= type-summand { type-summand }*

    type-summand ::= | constructor [slack-type-descriptor]

    constructor ::= simple-name

Sample type-sum:

    | Point XYpos | Line XYpos * XYpos

Restriction. The constructors of a type-sum must all be different simple-names.

The ordering of the type-summands has no significance: | Zero | Succ Peano denotes the same "sum type" as | Succ Peano | Zero.

A type-sum denotes a *sum type*, which is a type that is inhabited by "tagged values". A tagged value is a pair (C, v), in which C is a constructor and v is a typed value.

A type-sum introduces a number of embedders, one for each type-summand. In the discussion, we omit the optional embed keyword of the embedders. The embedders are similar to ops, and are explained as if they were ops, but note the Restriction specified under *Structors*.

For a type-sum T with type-summand C S, in which C is a constructor and S a type-descriptor, the corresponding pseudo-op introduced is typed as follows:

    op C : S -> T

It maps a value v of type S to the tagged value (C, v). If the type-summand is a single *parameter-less* constructor (the slack-type-descriptor is missing), the pseudo-op introduced is typed as follows:

    op C : T

It denotes the tagged value $(C, ())$, in which () is the inhabitant of the unit type (see under *Type-records*).

The sum type denoted by the type-sum then consists of the union of the ranges (for parameter-less constructors the values) of the pseudo-ops for all constructors.

The embedders are individually, jointly and severally *injective*, and jointly *surjective*.

This means, first, that for any pair of constructors $C1$ and $C2$ of *any* type-sum, and for any pair of values $v1$ and $v2$ of the appropriate type (to be omitted for

parameter-less constructors), the value of `C1 v1` is only equal to `C2 v2` when `C1` and `C2` are the same constructor of the *same* sum type, and `v1` and `v2` (which then are either both absent, or else must have the same type) are both absent or are the same value. In other words, whenever the constructors are different, or are from different type-sums, or the values are different, the results are different. (The fact that synonymous constructors of different types yield different values already follows from the fact that values in the models are typed.)

Secondly, for any value `u` of any sum type, there is a constructor `C` of that sum type and a value `v` of the appropriate type (to be omitted for parameter-less constructors), such that the value of `C v` is `u`. In other words, all values of a sum type can be constructed with an embedder.

For example, consider

```
type Peano =
   | Zero
   | Succ Peano

type Unique =
   | Zero
```

This means that there is a value `Zero` of type `Peano`, and further a function `Succ` that maps values of type `Peano` to type `Peano`. Then `Zero` and `Succ n` are guaranteed to be different, and each value of type `Peano` is either `Zero : Peano`, or expressible in the form `Succ (n : Peano)` for a suitable expression n. The expressions `Zero : Peano` and `Zero : Unique` denote different, entirely unrelated, values. (Note that `Unique` is *not* a subtype of `Peano`. Subtypes of a type can only be made with a type-restriction, for instance as in `(Peano | embed? Zero)`.) For recursively defined type-sums, see also the discussion under *Type-definitions*.

Note. Although the sum types `| Mono` and `| Mono ()` have exactly the same set of inhabitants when considered as untyped values, these two types are different, and the pseudo-ops they introduce have different types, only the second of which is a function type:

```
Mono : | Mono

Mono : () -> | Mono ()
```

2.5.2. Type-arrows

type-arrow ::= arrow-source **->** type-descriptor

arrow-source ::= type-sum | slack-type-descriptor

Sample type-arrow:

```
(a -> b) * b -> List a -> List b
```

In this example, the arrow-source is `(a -> b) * b`, and the (target) type-descriptor `List a -> List b`.

The *function type* `S -> T` is inhabited by precisely all *partial or total* functions from S to T. That is, function f has type `S -> T` if, and only if, for each value x of type S such that the value of $f\ x$ is defined, that value has type T. Functions can be constructed with lambda-forms, and be used in applications.

In considering whether two functions (of the same type) are equal, only the meaning on the domain type is relevant. Whether a function is undefined outside its domain type, or might return some value of some type, is immaterial to the semantics of Metaslang. (For a type-correct spec, the difference is unobservable.)

2.5.3. Type-products

type-product ::= tight-type-descriptor ***** tight-type-descriptor { ***** tight-type-descriptor }*

Sample type-product:

```
(a -> b) * b * List a
```

Note that a type-product contains at least two constituent tight-type-descriptors.

A type-product denotes a *product type* that has at least two "component types", represented by its tight-type-descriptors. The ordering of the component types is significant: unless S and T are the same type, the product type `S * T` is different from the type `T * S`. Further, the three types `(S * T) * U`, `S * (T * U)` and `S * T * U` are all different; the first two have two component types, while the last one has three. The inhabitants of the product type $T_1 * T_2 * ... * T_n$ are precisely all n-tuples (v_1, v_2, ... , v_n), where each v_i has type T_i, for $i = 1, 2, ... , n$. Values of a product type can be

constructed with tuple-displays, and component values can be extracted with tuple-patterns as well as with projectors.

2.5.4. Type-instantiations

type-instantiation ::= type-name actual-type-parameters

actual-type-parameters ::= closed-type-descriptor | (proper-type-list)

proper-type-list ::= type-descriptor , type-descriptor { , type-descriptor }*

Sample type-instantiation:

```
    Map (Nat, Boolean)
```

Restriction. The type-name must have been declared or defined as a parameterized type (see *Type-declarations*), and the number of type-descriptors in the actual-type-parameters must match the number of local-type-variables in the formal-type-parameters of the type-declaration and/or type-definition.

The type-descriptor represented by a type-instantiation is the type assigned for the combination of types of the actual-type-parameters in the indexed family of types for the type-name of the type-instantiation.

2.5.5. Type-names

type-name ::= name

Sample type-names:

```
    Key
    Calendar.Date
```

Restriction. At the spec level, a type-name may only be used if there is a type-declaration and/or type-definition for it in the current spec or in some spec that is imported (directly or indirectly) in the current spec. If there is a unique qualified-name for a given unqualified ending, the qualification may be omitted for a type-name used as a type-descriptor.

The type of a type-name is the type assigned to it in the model. (In this case, the context can not have superseded the original assignment.)

2.5.6. Type-records

type-record ::= { [field-typer-list] } | ( )

field-typer-list ::= field-typer { , field-typer }*

field-typer ::= field-name : type-descriptor

field-name ::= simple-name

Sample type-record:

    {center : XYpos, radius : Length}

Restriction. The field-names of a type-record must all be different.

Note that a type-record contains either no constituent field-typers, or else at least two.

A type-record is like a type-product, except that the components, called "fields", are identified by name instead of by position. The ordering of the field-typers has no significance: {center : XYpos, radius : Length} denotes the same *record type* as {radius : Length, center : XYpos}. Therefore we assume in the following, without loss of generality, that the fields are ordered lexicographically according to their field-names (as in a dictionary: a comes before ab comes before b) using some fixed collating order for all marks that may comprise a name. Then each field of a record type with n fields has a *position* in the range 1 to n. The inhabitants of the record type $\{F_1 : T_1, F_2 : T_2, ... , F_n : T_n\}$ are precisely all n-tuples $(v_1, v_2, ... , v_n)$, where each v_i has type T_i, for $i = 1, 2, ... , n$. The field-names of that record type are the field-names $F_1, ... , F_n$, and, given the lexicographic ordering, field-name F_i *selects* position i, for $i = 1, 2, ... , n$. Values of a record type can be constructed with record-displays, and field values can be extracted with record-patterns and (as for product types) with projectors.

For the type-record {}, which may be equivalently written as ( ), the record type it denotes has zero components, and therefore no field-names. This zero-component type has precisely one inhabitant, and is called the *unit type*. The unit type may equally well be considered a product type, and is the only type that is both a product and a record type.

2.5.7. Type-restrictions

type-restriction ::= **(** slack-type-descriptor **|** expression **)**

Sample type-restriction:

```
(Nat | even)
```

Restriction. In a type-restriction `(T | P)`, the expression `P` must be a predicate on the type `T`, that is, `P` must be a function of type `T -> Boolean`.

Note that the parentheses in `(T | P)` are mandatory.

The inhabitants of type `(T | P)` are precisely the inhabitants of type `T` that satisfy the predicate `P`, that is, they are those values `v` for which the value of `P v` is `true`.

If `P1` and `P2` are the same function, then `(T | P1)` and `(T | P2)` are equivalent, that is, they denote the same type. Furthermore, `(T | fn _ -> true)` is equivalent to `T`.

The type `(T | P)` is called a *subtype* of *supertype* `T`. Values can be shuttled between a subtype and its supertype and vice versa, in the direction from supertype to subtype only if the value satisfies predicate `P`. For example, in the expression `-1` the nat-literal `1` of type `Nat` is implicitly "coerced" to type `Integer` to accommodate the unary negation operator `-`, which has type `Integer -> Integer`.

Likewise, in the expression `7 div 2` the nat-literal `2` of type `Nat` is implicitly "coerced" to type `PosNat`, a subtype of `Nat`, to accommodate the division operator `div`, whose second argument has type `PosNat`. But note that this engenders the proof obligation that the value satisfies the predicate of the subtype.

These coercions extend to composed types. For example, an expression of type `List PosNat` may be used where a value of type `List Nat` is required. Conversely, an expression of type `List Nat` may be used in a context requiring `List PosNat` if the corresponding proof obligation can be discharged, namely that the value of the expression, in its context, satisfies the predicate `all posNat?` testing whether all elements of a list of naturals are positive.

2.5.8. Type-comprehensions

type-comprehension ::= **{** annotated-pattern **|** expression **}**

Sample type-comprehension:

```
{n : Nat | even n}
```

Restriction. In a type-comprehension {`P` : `T` | `E`}, the expression `E` must have type `Boolean`.

Type-comprehensions provide an alternative notation for type-restrictions that is akin to the common mathematical notation for set comprehensions. The meaning of type-comprehension {`P` : `T` | `E`} is the same as that of the type-restriction (`T` | fn `P` -> `E`). So the meaning of the example type-comprehension above is (`Nat` | fn n -> even n).

2.5.9. Type-quotients

type-quotient ::= closed-type-descriptor / closed-expression

Sample type-quotient:

```
Nat / (fn (m, n) -> m rem 3 = n rem 3)
```

Restriction. In a type-quotient `T` / `Q`, the expression `Q` must be a (binary) predicate on the type `T` * `T` that is an equivalence relation, as explained below.

Equivalence relation. Call two values x and y of type T "Q-related" if (x, y) satisfies Q. Then Q is an *equivalence relation* if, for all values x, y and z of type T, x is Q-related to itself, y is Q-related to x whenever x is Q-related to y, and x is Q-related to z whenever x is Q-related to y and y is Q-related to z. The equivalence relation Q then partitions the inhabitants of T into *equivalence classes*, being the maximal subsets of T containing mutually Q-related members. These equivalence classes will be called "Q-equivalence classes".

The inhabitants of the *quotient type* `T` / `Q` are precisely the Q-equivalence classes into which the inhabitants of T are partitioned by Q. For the example above, there are three equivalence classes of natural numbers leaving the same remainder on division by 3: the sets {0, 3, 6, ...}, {1, 4, 7, ...} and {2, 5, 8, ...}, and so the quotient type has three inhabitants.

2.6. Expressions

```
expression ::=
        lambda-form
      | case-expression
      | let-expression
      | if-expression
      | quantification
      | unique-solution
      | annotated-expression
      | tight-expression

tight-expression ::=
        application
      | closed-expression

closed-expression ::=
        op-name
      | local-variable
      | literal
      | field-selection
      | tuple-display
      | record-display
      | sequential-expression
      | list-display
      | monadic-expression
      | structor
      | ( expression )
      | ( inbuilt-op )
```

inbuilt-op ::= inbuilt-prefix-op | inbuilt-infix-op

inbuilt-prefix-op ::= **~**

inbuilt-infix-op ::= **<=>** | **=>** | **||** | **&&** | **=** | **~=** | **<<**

(The distinctions tight- and closed- for expressions lack semantic significance, and merely serve the purpose of avoiding grammatical ambiguities.)

Sample expressions:

```
fn (s : String) -> s ^ "."
case z of {re = x, im = y} -> {re = x, im = -y}
let x = x + 1 in f(x, x)
```

```
if x <= y then x else y
fa(x,y) (x <= y)  <=>  ((x < y) or (x = y))
f(x, x)
[] : List Arg
abs(x-y)
++
x
3260
z.re
("George", Poodle : Dog, 10)
{name = "George", kind = Poodle : Dog, age = 10}
(writeLine "key not found"; embed Missing)
["Sun", "Mon", "Tue", "Wed", "Thu", "Fri", "Sat"]
project 2
(n + 1)
(||)
```

Restriction. Like all polymorphic or type-ambiguous constructs, an expression can only be used in a context if its type can be inferred uniquely, given the expression and the context. This restriction will not be repeated for the various kinds of expressions defined in the following subsections.

The meaning of a parenthesized expression `(E)` is the same as that of the enclosed expression `E`. The meaning of the parenthesized inbuilt-prefix-op `(P)` is the same as that of the lambda-form `fn x -> P x`. The meaning of a parenthesized inbuilt-infix-op `(I)` is the same as that of the lambda-form `fn (x,y) -> x I y`. Note that this function is strict in both arguments, unlike `I` itself.

The various other kinds of expressions not defined here are described each in their following respective sections, with the exception of local-variable, whose meaning as an expression is described below.

Restriction. A local-variable may only be used as an expression if it occurs in the scope of the local-variable-list of a quantification or of a variable-pattern in which it is introduced.

Disambiguation. A single simple-name used as an expression is a local-variable when it occurs in the scope of a local-variable-list or variable-pattern in which a synonymous local-variable is introduced, and then it identifies the textually most recent introduction. Otherwise, the simple-name is an op-name or an embedder; for the disambiguation between the latter two, see *Embedders*.

A local-variable used as an expression has the typed value assigned to it in the environment.

2.6.1. Lambda-forms

lambda-form ::= **fn** match

Sample lambda-forms:

```
fn (s : String) -> s ^ "."

fn | Point _       -> 0
   | Line(z0, z1) -> dist(z0, z1)
```

The value of a lambda-form is a partial or total function. If the value determined for a lambda-form as described below is not a total function, the context must enforce that the function can not be applied to values for which it is undefined. Otherwise, the spec is ill formed. Specware 4.2 does not attempt to generate proof obligations for establishing this.

The type of a lambda-form is that of its match. The meaning of a given lambda-form of type $S \rightarrow T$ is the function f mapping each inhabitant x of S to a value y of type T, where y is the return value of x for the match of the lambda-form. If the match accepts each x of type S (for acceptance and return value, see the section on *Matches*) function f is total; otherwise it is partial, and undefined for those values x rejected.

In case of a recursive definition, the above procedure may fail to determine a value for y, in which case function f is not total, but undefined for x.

2.6.2. Case-expressions

case-expression ::= **case** expression **of** match

Sample case-expressions:

```
case z of {re = x, im = y} -> {re = x, im = -y}

case s of
   | Empty -> true
   | Push {top = _, pop = rest} -> hasBottom? rest
```

The value of a case-expression `case E of M` is the same as that of the application `(fn M) (E)`.

2.6.3. Let-expressions

let-expression ::= **let** let-bindings **in** expression

let-bindings ::= recless-let-binding | rec-let-binding-sequence

recless-let-binding ::= pattern equals expression

rec-let-binding-sequence ::= rec-let-binding { rec-let-binding }*

rec-let-binding ::=
    def simple-name formal-parameter-sequence [**:** type-descriptor] equals expression

formal-parameter-sequence ::= formal-parameter { formal-parameter }*

Sample let-expressions:

```
let x = x + e in f(x, x)
let def f x = x + e in f (f x)
```

In the case of a recless-let-binding (recless = recursion-less), the value of the let-expression `let P = A in E` is the same as that of the application `(fn P -> E)(A)`. For the first example above, this amounts to `f(x + e, x + e)`. Note that `x = x + e` is not interpreted as a recursive definition.

In case of a rec-let-binding-sequence (rec = recursive), the rec-let-bindings have the role of "local" op-definitions; that is, they are treated exactly like op-definitions except that they are interpreted in the local environment instead of the global model. For the second example above, this amounts to `(x + e) + e`. (If `e` is a local-variable in this scope, the definition of `f` can not be "promoted" to an op-definition, which would be outside the scope binding `e`.) A spec with rec-let-bindings can be transformed into one without such by creating op-definitions for each rec-let-binding that take additional arguments, one for each of the local-variables referenced. For the example, in which `f` references local-variable `e`, the op-definition for the "extended" op f^+ would be `def f+ e x = x + e`, and the let-expression would become `f+ e (f+ e x)`. The only difference in meaning is that the models of the transformed spec assign a value to the newly introduced op f^+.

Note that the first occurrence of `x` in the above example of a rec-let-binding is a variable-pattern and the second-occurrence is in its scope; the third and last occurrence of `x`, however, is outside the scope of the first `x` and identifies an op or

local-variable x introduced elsewhere. So, without change in meaning, the rec-let-binding can be changed to:

```
let def f xena = xena + e in f (f x)
```

2.6.4. If-expressions

if-expression ::= **if** expression **then** expression **else** expression

Sample if-expression:

```
if x <= y then x else y
```

The value of an if-expression if *B* then *T* else *F* is the same as that of the case-expression case *B* of true -> (*T*) | false -> (*F*).

2.6.5. Quantifications

quantification ::= quantifier **(** local-variable-list **)** expression

quantifier ::= **fa** | **ex** | **ex1**

local-variable-list ::= annotable-variable { **,** annotable-variable }*

annotable-variable ::= local-variable [**:** type-descriptor]

local-variable ::= simple-name

Sample quantifications:

```
fa(x) norm (norm x) = norm x
ex(e : M) fa(x : M) x <*> e = x & e <*> x = x
```

Restriction. Each local-variable of the local-variable-list must be a different simple-name.

Quantifications are non-constructive, even when the domain type is finitely enumerable. The main uses are in type-restrictions and type-comprehensions, and

claims. The type of a quantification is `Boolean`. There are three kinds of quantifiers: `fa`, for "universal quantifications" --- fa = for all; `ex`, for "existential quantifications" --- ex = there exists; and `ex1`, for "uniquely existential quantifications" --- ex1 = there exists one.

The value of a quantification `fa` `(V)` `E`, in which V is a local-variable-list and E is an expression, is determined as follows. Let M be the match `(V)` `->` `E`. If M has return value `true` for each value x in its domain, the value of the quantification is `true`; otherwise it is `false`.

The value of a quantification `ex` `(V))` `E` in which V is a local-variable-list and E is an expression, is determined as follows. Let M be the match `(V)` `->` `E`. If M has return value `true` for at least one value x in its domain, the value of the quantification is `true`; otherwise it is `false`.

The value of a quantification `ex1` `(V))` `E` in which V is a local-variable-list and E is an expression, is determined as follows. Let M be the match `(V)` `->` `E`. If M has return value `true` for precisely one value x in its domain, the value of the quantification is `true`; otherwise it is `false`.

Note that a quantifier must be followed by an opening parenthesis `(` . So `fa x (x =` `x)`, for example, is ungrammatical.

2.6.6. Unique-solutions

unique-solution `::=` **the (** local-variable-list **)** expression

Sample unique-solution:

```
the(x : S) f(x) = y
```

Restriction, Each local-variable of the local-variable-list must be a different simple-name.

Restriction. The type of the expression must be `Boolean`.

Restriction. A unique-solution `the` `(V)` `E` may only be used in a context where the value of `ex1` `(V)` `E` is `true`.

Unique-solutions are non-constructive, even when the domain type is finitely enumerable. The type of a unique-solution is the type of its local-variable-list.

The value of a unique-solution `the` `(V)` `E`, in which V is a local-variable-list and E is an expression, is determined as follows. Let M be the match `(V)` `->` `E`. The value

of the unique-solution is then the unique value `x` in the domain of `M` such that the match `(V) -> E` has return value `true` for `x`.

2.6.7. Annotated-expressions

annotated-expression `::=` tight-expression **:** type-descriptor

Restriction. In an annotated-expression `E : T`, the expression `E` must have type `T`.

Sample annotated-expression:

```
[] : List Arg
Positive : Sign
```

The value of an annotated-expression `E : T` is the value of `E`.

The type of some expressions is polymorphic. For example, for any type `T`, `[]` denotes the empty list of type `List T`. Likewise, constructors of parameterized sum types can be polymorphic, as the constructor `None` of

```
type Option a = | Some a | None
```

Further, overloaded constructors have an ambiguous type. By annotating such polymorphic or type-ambiguous expressions with a type-descriptor, their type can be disambiguated, which is required unless an unambiguous type can already be inferred from the context. Annotation, even when redundant, can further help to increase clarity.

2.6.8. Applications

application `::=` prefix-application | infix-application

prefix-application `::=` application-head actual-parameter

application-head `::=`
        closed-expression
    | inbuilt-prefix-op
    | prefix-application

actual-parameter `::=` closed-expression

```
infix-application  ::=  operand  infix-operator  operand
```

```
operand  ::=  tight-expression
```

```
infix-operator  ::=  op-name  |  inbuilt-infix-op
```

Sample applications:

```
f (x, x)
f x (g y)
x + 1
```

Restriction. An infix-operator, whether qualified or unqualified, can not be used without more as an actual-parameter or operand (and in the case of an inbuilt-op, it can not be used without more as any other kind of expression either). To use an infix-operator in such cases, it must be enclosed in parentheses, as for example in the prefix-applications `foldl (+) 0` and `foldl ( *) 1` or the infix-application `(<) o ival`. Note the space between "`(`" and "`*`", since without space "`(*`" signals the start of a comment.

Restriction. An op-name can be used as an infix-operator only if it has been declared as such in an op-declaration (see under *Op-declarations*).

Disambiguation. An infix-application $P\ M\ Q\ N\ R$, in which P, Q and R are operands and M and N are infix-operators, is interpreted as either $(P\ M\ Q)\ N\ R$ or $P\ M\ (Q\ N\ R)$. The choice is made as follows. If M has higher priority than N, or the priorities are the same but M is left-associative, the interpretation is $(P\ M\ Q)\ N\ R$. In all other cases the interpretation is $P\ M\ (Q\ N\ R)$. For example, given

```
op @ infixl 10: Nat * Nat -> Nat
op $ infixr 20: Nat * Nat -> Nat
```

the following interpretations hold:

```
1 $ 2 @ 3  =  (1 $  2) @ 3
1 @ 2 @ 3  =  (1 @  2) @ 3
1 @ 2 $ 3  =   1 @ (2  $ 3)
1 $ 2 $ 3  =   1 $ (2  $ 3)
```

Note that no type information is used in the disambiguation. If `(1 @ 2) $ 3` is type-correct but `1 @ (2 $ 3)` is not, the formula `1 @ 2 $ 3` is type-incorrect, since its interpretation is.

For the application of this disambiguation rule, the inbuilt-ops have fixity as suggested by the following pseudo-op-declarations:

```
op <=>  infixr 12 : Boolean * Boolean -> Boolean
op =>   infixr 13 : Boolean * Boolean -> Boolean
op ||   infixr 14 : Boolean * Boolean -> Boolean
op &&   infixr 15 : Boolean * Boolean -> Boolean
op =    infixr 20 : [a]  a * a        -> Boolean
op ~=   infixr 20 : [a]  a * a        -> Boolean
op <<   infixl 25 : {x:A, ... , y:B, ...} * {x:A, ... , z:C, ...}
                 -> {x:A, ... , y:B, ... , z:C, ...}
```

Restriction. In an application $H\ P$, in which H is an application-head and P an actual-parameter, the type of (H) must be some function type $S\ \text{->}\ T$, and then P must have the domain type S. The type of the whole application is then T. In particular, in an application $\sim P$ the type of both P and the application is `Boolean`.

The meaning of prefix-application $\sim P$ is the same as that of the if-expression `if P then false else true`.

The value of prefix-application $H\ P$, in which application-head H is a closed-expression or another prefix-application, is the value returned by function (H) for the argument value P.

The meaning of infix-application $P\ N\ Q$, in which P and Q are operands and N is an op-name, is the same as that of the prefix-application $N(P,\ Q)$.

The meaning of infix-application $P\ \text{=>}\ Q$, in which P and Q are operands, is the same as that of the if-expression `if P then Q else true`.

The meaning of infix-application $P\ ||\ Q$, in which P and Q are operands, is the same as that of the if-expression `if P then true else Q`.

The meaning of infix-application $P\ \&\&\ Q$, in which P and Q are operands, is the same as that of the if-expression `if P then Q else false`.

The value of infix-application $P\ \text{=}\ Q$, in which P and Q are operands, is `true` if P and Q have the same value, and `false` otherwise. P and Q must have the same type, or else have types that are subtypes of the same supertype. In the latter case, the comparison is the same as for the values of the operands coerced to the supertype, so, for example, the value of `(1:Nat)=(1:PosNat)` is `true`.

The meaning of infix-application $P\ \sim=\ Q$, in which P and Q are operands, is the same as that of the prefix-application $\sim(P\ =\ Q)$.

An infix-application $P\ \text{<<}\ Q$ is also called a "record update". In a record update $P\ \text{<<}\ Q$, in which P and Q are operands, P and Q must have record types, referred to as S and T, respectively. Moreover, for each field-name F these types S and T have in

common, the field types for F in S and T must be the same, or be subtypes of the same supertype. The type of P `<<` Q is then the record type R whose field-names are formed by the union of the field-names of S and T, where for each field-name F in that union, the type of field F in R is that of field F in T if F is a field of T, and otherwise the type of field F in S. Likewise, the value of P `<<` Q is the record value of type R whose field value of each field F is that of field F in Q if F is a field of T, and otherwise the field value of field F in P. So, for example, the value of {a=1, b=#z} `<<` {a=2, c=true} is {a=2, b=#z, c=true}: fields of the right-hand side operand take precedence over the left-hand side when present in both.

2.6.9. Op-names

```
op-name ::= name
```

Sample op-names:

```
length
>=
DB_LOOKUP.Lookup
```

Restriction. An op-name may only be used if there is an op-declaration and/or op-definition for it in the current spec or in some spec that is imported (directly or indirectly) in the current spec. If there is a unique qualified-name for a given unqualified ending that is type-correct in the context, the qualification may be omitted for an op-name used as an expression. So overloaded ops may only be used as such when their type can be disambiguated in the context.

The value of an op-name is the value assigned to it in the model. (In this case, the context can not have superseded the original assignment.)

2.6.10. Literals

```
literal ::=
        boolean-literal
      | nat-literal
      | char-literal
      | string-literal
```

Sample literals:

```
true
3260
#z
"On/Off switch"
```

Restriction: No whitespace is allowed anywhere inside any kind of literal, except for "significant" whitespace in string-literals, as explained there.

Literals provide denotations for the inhabitants of the inbuilt and "base-library" types `Boolean`, `Nat`, `Char` and `String`. The value of a literal is independent of the environment.

(There are no literals for the base-library type `Integer`. For nonnegative integers, a nat-literal can be used. For negative integers, apply the unary base-library op `-`, which negates an integer: `-1` denote the negative integer –1.)

2.6.10.1. Boolean-literals

boolean-literal ::= **true** | **false**

Sample boolean-literals:

```
true
false
```

The type `Boolean` has precisely two inhabitants, the values of `true` and `false`.

Note that `true` and `false` are not constructors. So `embed true` is ungrammatical.

2.6.10.2. Nat-literals

nat-literal ::= decimal-digit { decimal-digit } *

Sample nat-literals:

```
3260
007
```

The type-descriptor `Nat` is, by definition, the subtype of `Integer` restricted to the nonnegative integers 0, 1, 2, ... , which we identify with the natural numbers. The value of a nat-literal is the natural number of which it is a decimal representation; for

example, the nat-literal 3260 denotes the natural number 3260. Leading decimal-digits 0 have no significance: both 007 and 7 denote the number 7.

2.6.10.3. Char-literals

char-literal ::= #char-literal-glyph

char-literal-glyph ::= char-glyph | "

char-glyph ::=
    letter
  | decimal-digit
  | other-char-glyph

other-char-glyph ::=
    ! | : | @ | # | $ | % | ^ | & | * | (|) | _ | - | + | =
  | | | ~ | ` | . | , | < | > | ? | / | ; | ' | [|] | { | }
  | \\ | \"
  | \a | \b | \t | \n | \v | \f | \r | \s
  | \x hexadecimal-digit hexadecimal-digit

hexadecimal-digit ::=
    decimal-digit
  | a | b | c | d | e | f
  | A | B | C | D | E | F

Sample char-literals:

```
#z
#\x7a
```

The type Char is inhabited by the 256 8-bit *characters* occupying decimal positions 0 through 255 (hexadecimal positions 00 through FF) in the ISO 8859-1 code table. The first 128 characters of that code table are the traditional ASCII characters (ISO 646). (Depending on the operating environment, in particular the second set of 128 characters -- those with "the high bit set" -- may print or otherwise be visually presented differently than intended by the ISO 8859-1 code.) The value of a char-literal is a character of type Char.

The value of a char-literal #G, where G is a char-glyph, is the character denoted by G. For example, #z is the character that prints as z. The two-mark char-literal #" provides

a variant notation of the three-mark char-literal `#\"` and yields the character `"` (decimal position 34).

Each one-mark char-glyph C denotes the character that "prints" as C. The two-mark char-glyph `\\` denotes the character `\` (decimal position 92), and the two-mark char-glyph `\"` denotes the character `"` (decimal position 34).

Notations are provided for denoting eight "non-printing" characters, which, with the exception of the first, are meant to regulate lay-out in printing; the actual effect may depend on the operating environment:

glyph	decimal	name
`\a`	7	bell
`\b`	8	backspace
`\t`	9	horizontal tab
`\n`	10	newline
`\v`	11	vertical tab
`\f`	12	form feed
`\r`	13	return
`\s`	32	space

Finally, every character can be obtained using the hexadecimal representation of its position. The four-mark char-glyph `\x`$H_1 H_0$ denotes the character with hexadecimal position $H_1 H_0$, which is decimal position 16 times the decimal value of hexadecimal-digit H_1 plus the decimal value of hexadecimal-digit H_0, where the decimal value of the digits `0` through `9` is conventional, while the six extra digits `A` through `F` correspond to 10 through 15. The case (lower or upper) of the six extra digits is not significant. For example, `\x7A` or equivalently `\x7a` has decimal position 16 times 7 plus 10 = 122, and either version denotes the character `z`. The "null" character can be obtained by using `\x00`.

2.6.10.4. String-literals

string-literal ::= `"` string-body `"`

string-body ::= { string-literal-glyph }*

string-literal-glyph ::= char-glyph | significant-whitespace

```
significant-whitespace ::= space | tab | newline
```

The presentation of a significant-whitespace is the whitespace suggested by the name (space, tab or newline).

Sample string-literals:

```
""
"see page"
"see\spage"
"the symbol ’ is a single quote"
"the symbol \" is a double quote"
```

The type `String` is inhabited by the *strings*, which are (possibly empty) sequences of characters. The type `String` is primitive; it is a different type than the isomorphic type `List Char`, and the list operations can not be directly applied to strings.

The value of a string-literal is the sequence of characters denoted by the string-literal-glyphs comprising its string-body, where the value of a significant-whitespace is the whitespace character suggested by the name (space, horizontal tab or newline). For example, the string-literal `"seepage"` is different from `"see page"`; the latter denotes an eight-character string of which the fourth character is a space. The space can be made explicit by using the char-glyph `\s`.

When a double-quote character `"` is needed in a string, it must be escaped, as in `"[6’2\"]"`, which would print like this: `[6’2"]`.

2.6.11. Field-selections

```
field-selection ::= closed-expression . field-selector
```

```
field-selector ::= nat-literal | field-name
```

Disambiguation. A closed-expression of the form `M.N`, in which `M` and `N` are simple-names, is interpreted as an op if `M.N` occurs as the op-name of an op-declaration or op-definition in the spec in which it occurs or in the set of simple-names imported from another spec through an import-declaration. Otherwise, `M.N` is interpreted as a field-selection. (The effect of a field-selection can always be obtained with a projector.)

Sample field-selections:

```
triple.2
z.re
```

A field-selection `E.F` is a convenient notation for the equivalent expression `(project F E)`. (See under *Projectors*.)

2.6.12. Tuple-displays

tuple-display `::=` **(** tuple-display-body **)**

tuple-display-body `::=` [expression **,** expression { **,** expression }*]

Sample tuple-display:

```
("George", Poodle : Dog, 10)
```

Note that a tuple-display-body contains either no expressions, or else at least two.

The value of a tuple-display whose tuple-display-body is not empty, is the tuple whose components are the respective values of the expressions of the tuple-display-body, taken in textual order. The type of that tuple is the "product" of the corresponding types of the components. The value of `()` is the empty tuple, which is the sole inhabitant of the unit type `()`. (The fact that the notation `()` does double duty, for a type-descriptor and as an expression, creates no ambiguity. Note also that -- unlike the empty list-display `[]` -- the expression `()` is monomorphic, so there is no need to ever annotate it with a type-descriptor.)

2.6.13. Record-displays

record-display `::=` **{** record-display-body **}**

record-display-body `::=` [field-filler { **,** field-filler }*]

field-filler `::=` field-name equals expression

Sample record-display:

```
{name = "George", kind = Poodle : Dog, age = 10}
```

The value of a record-display is the record whose components are the respective values of the expressions of the record-display-body, taken in the lexicographic order of the field-names, as discussed under *Type-records*. The type of that record is the record type with the same set of field-names, where the type for each field-name *F* is the type of the corresponding type of the component selected by *F* in the record. The value of {} is the empty tuple, which is the sole inhabitant of the unit type (). (For expressions as well as for type-descriptors, the notations {} and () are fully interchangeable.)

2.6.14. Sequential-expressions

sequential-expression ::= **(** open-sequential-expression **)**

open-sequential-expression ::= void-expression **;** sequential-tail

void-expression ::= expression

sequential-tail ::= expression | open-sequential-expression

Sample sequential-expression:

```
(writeLine "key not found"; embed Missing)
```

A sequential-expression (*V*; *T*) is equivalent to the let-expression `let _ = V` `in (T)`. So the value of a sequential-expression (*V*$_1$; ... ; *V*$_n$; *E*) is the value of its last constituent expression *E*.

Sequential-expressions can be used to achieve non-functional "side effects", effectuated by the elaboration of the void-expressions, in particular the output of a message. This is useful for tracing the execution of generated code. The equivalent effect of the example above can be achieved by a let-binding:

```
let _ = writeLine "key not found" in
embed Missing
```

(If the intent is to temporarily add, and later remove or disable the tracing output, this is probably a more convenient style, as the modifications needed concern a single full text line.) Any values resulting from elaborating the void-expressions are discarded.

2.6.15. List-displays

list-display ::= **[** list-display-body **]**

list-display-body ::= [expression { **,** expression }*]

Sample list-display:

```
["Sun", "Mon", "Tue", "Wed", "Thu", "Fri", "Sat"]
```

Restriction. All expressions of the list-display-body must have the same type.

Note that a list-display `[]` with empty list-display-body is polymorphic, and may need to be type-disambiguated, for example with a type annotation. In a case like `[[],` `[1]]`, there is no need to disambiguate `[]`, since the above restriction already implies that `[]` here has the same type as `[1]`, which has type `List Nat`.

The parameterized type `List`, although a base-library type, is actually not primitive, but defined by:

```
type List a =
  | Nil
  | Cons a * List a
```

The empty list-display `[]` denotes the same list as the expression `Nil`, a singleton list-display `[E]` denotes the same list as the expression `Cons (E, Nil)`, and a multi-element list-display `[E_1, E_2, ... , E_n]` denotes the same list as the expression `Cons (E_1, [E_2, ... , E_n])`.

2.6.16. Monadic-expressions

monadic-expression ::= **{** open-monadic-expression **}**

open-monadic-expression ::= monadic-statement **;** monadic-tail

monadic-statement ::= expression | monadic-binding

monadic-binding ::= pattern **<-** expression

monadic-tail ::= expression | open-monadic-expression

Sample monadic-expression:

```
{x <- a; y <- b; f(x, y)}
```

Restriction. Monadic-expressions can only be used in a context containing the following spec, or a refinement thereof, possibly qualified, as a sub-spec (see under *Substitutions*):

```
spec
  type Monad a

  op monadBind : [a,b] (Monad a) * (a -> Monad b) -> Monad b
  op monadSeq  : [a,b] (Monad a) *      (Monad b) -> Monad b
  op return : [a] a -> Monad a

  axiom left_unit is
    [a,b] fa (f : a -> Monad b, x : a)
      monadBind (return x, f) = f x

  axiom right_unit is
    [a] fa (m : Monad a)
      monadBind (m, return) = m

  axiom associativity is
    [a,b,c] fa (m : Monad a, f : a -> Monad b, h : b->Monad c)
      monadBind (m, (fn x -> monadBind (f x, h))) =
      monadBind (monadBind (m, f), h)

  axiom non_binding_sequence is
    [a] fa (f : Monad a, g : Monad a)
      monadSeq (f, g) = monadBind (f, fn _ -> g)

endspec
```

(This spec can be found, qualified with `Monad`, in the library spec `/Library/Structures/Data/Monad`.) A monadic-expression may further only be used when the non-monadic expression it is equivalent to (see below) is itself a valid expression.

A monadic-expression {*M*} is equivalent to the open-monadic-expression *M*.

A monadic-tail *E*, where *E* is an expression, is equivalent to the expression *E*.

A monadic-tail *M*, where *M* is an open-monadic-expression, is equivalent to the open-monadic-expression *M*.

An open-monadic-expression `E; T`, where `E` is an expression, is equivalent to the application `monadSeq (E, T')`, where `T'` is an expression that is equivalent to the monadic-tail `T`.

An open-monadic-expression `P <- E; T` is equivalent to the application `monadBind (E, fn P -> T')`, where `T'` is an expression that is equivalent to the monadic-tail `T`.

2.6.17. Structors

```
structor ::=
        projector
      | quotienter
      | chooser
      | embedder
      | embedding-test
```

The structors are a medley of constructs, all having polymorphic or type-ambiguous function types and denoting special functions that go between structurally related types, such as the constructors of sum types and the destructors of product types.

Restriction. Like all polymorphic or type-ambiguous constructs, a structor can only be used in a context where its type can be inferred uniquely. This restriction will not be repeated for the various kinds of structors described in the following subsections.

For example, the following well-formed spec becomes ill formed when any of the type annotations is omitted:

```
spec
  def [a] p2 = project 2 : String *  a  ->  a
  def     q2 = project 2 : String * Nat -> Nat
endspec
```

2.6.17.1. Projectors

projector ::= **project** field-selector

Sample projectors:

```
project 2
project re
```

When the field-selector is some nat-literal with value i, it is required that i be at least 1. The type of the projector is a function type (whose domain type is a product type) of the form $T_1 * T_2 * \ldots * T_n \rightarrow T_i$, where n is at least i, and the value of the projector is the function that maps each n-tuple $(v_1, v_2, \ldots, v_n)$ inhabiting the domain type to its ith component v_i.

When the field-selector is some field-name F, the type of the projector is a function type (whose domain type is a record type) of the form $\{F_1 : T_1, F_2 : T_2, \ldots, F_n : T_n\} \rightarrow T_i$, where F is the same field-name as F_i for some natural number i in the range 1 through n. Assuming that the fields are lexicographically ordered by field-name (see under *Type-records*), the value of the projector is the function that maps each n-tuple $(v_1, v_2, \ldots, v_n)$ inhabiting the domain type to its ith component v_i.

2.6.17.2. Quotienters

quotienter ::= **quotient [** type-name **]**

Sample quotienter:

```
quotient[Q]
```

Restriction. In a quotienter `quotient[Q]`, Q must be defined as a *quotient type*. e.g. $Q = T / q$.

The type of quotienter `quotient[Q]`, where $Q = T / q$, is the function type $T \rightarrow Q$, that is, it goes from the base type to the quotient type. The value of the quotienter is the function that maps each inhabitant of type T to the q-equivalence class inhabiting Q of which it is a member.

For example, given

```
def congMod3 : Nat * Nat -> Boolean =
  (fn (m, n) -> m rem 3 = n rem 3)

type Z3 = Nat / congMod3
```

we have the typing

```
quotient[Z3] : Nat -> Z3
```

and the function maps, for example, the number 5 to the equivalence class {2, 5, 8, ...}, which is one of the three inhabitants of `Z3`.

2.6.17.3. Choosers

chooser ::= **choose [** type-name **]**

Sample chooser:

```
choose[Q]
```

Restriction. In a chooser `choose[Q]`, Q must be defined as a *quotient type*.

The type of chooser `choose[Q]`, where `Q = T / q`, is a function type of the form F `-> (Q -> R)`, where F is the subtype of `T -> R` consisting of the q-constant (explained below) functions. Expressed more formally, F is the type `{f : T -> R | fa((x,y) : T * T) q(x,y) => f x = f y}`, where the simple-names f, x and y must be replaced by "fresh" simple-names not clashing with simple-names already in use in T, R or q.

The value of the chooser is the function mapping each q-constant (explained below) function f inhabiting type `T -> R` to the function of type `Q -> R` that maps each inhabitant C of Q to `f x`, where x is any member of C. Expressed symbolically, using a pseudo-function `any` that arbitrarily picks any member from a nonempty set, this is the function

```
fn f  ->  fn C -> f (any C)
```

The requirement of q-constancy on f is precisely what is needed to make this function insensitive to the choice made by `any`.

Function f is q-constant if, for each q-equivalence class C inhabiting Q, `f x` equals `f y` for any two values x and y that are members of C, or f is undefined on all members of C. (Since the result of f is constant across each equivalence class, it does not matter which of its elements is selected by `any`.) For example -- continuing the example of the previous section -- function `fn n -> n*n rem 3` is congMod3-constant; for the equivalence class {2, 5, 8, ...}, for example, it maps each member to the same value 1. So `choose[Z3]` maps the inhabitant {2, 5, 8, ...} of type `Z3` to the natural number 1.

The most discriminating q-constant functions, where q has type `T * T -> Boolean`, are quotienters of the form `quotient[Q]`, where `Q = T / q`. Since all such Q are

isomorphic to one another (but not necessarily equal), all these quotinters are isomorphic to one another.

The meaning of `choose[Q] (fn x -> E) A` is the same as that of the let-expression `let quotient[Q] x = A in E`. Indeed, often a quotient-pattern offers a more convenient way of expressing the intention of a chooser. Note, however, the remarks on the proof obligations for quotient-patterns.

2.6.17.4. Embedders

embedder ::= [**embed**] constructor

Sample embedders:

```
Nil
embed Nil
Cons
embed Cons
```

Disambiguation. If an expression consists of a single simple-name, which, in the context, is both the simple-name of a constructor and the simple-name of an op or a local-variable in scope, then it is interpreted as the latter of the various possibilities. For example, in the context of

```
type Answer = | yes | no

def yes = no : Answer

def which (a : Answer) = case a of
  | yes -> "Yes!"
  | no  -> "Oh, no!"
```

the value of `which yes` is `"Oh, no!"`, since `yes` here is disambiguated as identifying the op `yes`, which has value `no`. The interpretation as embedder is forced by using the `embed` keyword: the value of `which embed yes` is `"Yes!"`. By using simple-names that begin with a capital letter for constructors, and simple-names that do not begin with a capital letter for ops and local-variables, the risk of an accidental wrong interpretation can be avoided.

The semantics of embedders is described in the section on *Type-sums*. The presence or absence of the keyword `embed` is not significant for the meaning of the construct (although it may be required for grammatical disambiguation, as described above).

2.6.17.5. Embedding-tests

embedding-test ::= **embed?** constructor

Sample embedding-test:

```
embed? Cons
```

Restriction. The type of an embedding-test `embed?` `C` must be of the form `T -> Boolean`, where `T` is a sum type that has a constructor `C`.

The value of embedding-test `embed?` `C` is the predicate that returns `true` if the argument value -- which, as inhabitant of a sum type, is tagged -- has tag `C`, and otherwise `false`. The embedding-test can be equivalently rewritten as

```
fn
  | C _   ->  true
  | _     ->  false
```

where the wildcard _ in the first branch is omitted when `C` is parameter-less.

In plain words, `embed?` `C` tests whether its sum-typed argument has been constructed with the constructor `C`. It is an error when `C` is not a constructor of the sum type.

2.7. Matches and Patterns

2.7.1. Matches

match ::= [**|**] branch { **|** branch }*

branch ::= pattern [guard] **->** expression

guard ::= **|** expression

Sample matches:

```
{re = x, im = y} -> {re = x, im = -y}

   Empty -> true
```

```
      | Push {top = _, pop = rest} -> hasBottom? rest


      | Empty -> true
      | Push {top = _, pop = rest} -> hasBottom? rest


      | Line(z0, z1) | z0 ~= z1 -> dist(z0, z1)
```

Restriction. In a match, given the environment, there must be a unique type S to which the pattern of each branch conforms, and a unique type T to which the expression of each branch conforms, and then the match has type $S \to T$. The pattern of each branch then has type S.

Restriction. The type of the expression of a guard must be `Boolean`

Disambiguation. If a branch could belong to several open matches, it is interpreted as being a branch of the textually most recently introduced match. For example,

```
      case x of
        | A -> a
        | B -> case y of
                 | C -> c
        | D -> d
```

is not interpreted as suggested by the indentation, but as

```
      case x of
        | A -> a
        | B -> (case y of
                 | C -> c
                 | D -> d)
```

If the other interpretation is intended, the expression introducing the inner match needs to be parenthesized:

```
      case x of
        | A -> a
        | B -> (case y of
                 | C -> c)
        | D -> d
```

Acceptance and return value y, if any, of a value x for a given match are determined as follows. If each branch of the match rejects x (see below), the whole match rejects

x, and does not return a value. Otherwise, let B stand for the textually first branch accepting x. Then y is the return value of x for B.

The meaning of a "guardless" branch P -> R, where P is a pattern and R an expression, is the same as that of the branch P | true -> R with a guard that always succeeds.

Acceptance and return value y, if any, of a value x for a branch P | G -> R in an environment C are determined as follows. If pattern P rejects x, the branch rejects x, and does not return a value. (For acceptance by a pattern, see under *Patterns*.) Otherwise, let C' be environment C extended with the acceptance binding of pattern P for x. If pattern P accepts x, but the value of expression G in the environment C' is false, the branch also rejects x, and does not return a value. Otherwise, when the pattern accepts x and the guard succeeds, the branch accepts x and the return value y is the value of expression R in the environment C'.

For example, in

```
case z of
    | (x, true)  -> Some x
    | (_, false) -> None
```

if z has value (3, true), the first branch accepts this value with acceptance binding x = 3. The value of Some x in the extended environment is then Some 3. If z has value (3, false), the second branch accepts this value with empty acceptance binding (empty since there are no "accepting" local-variables in pattern (_, false)), and the return value is None (interpreted in the original environment).

Here is a way of achieving the same result using a guard:

```
case z of
    | (x, b) | b -> Some x
    | _           -> None
```

2.7.2. Patterns

```
pattern ::=
      annotated-pattern
    | tight-pattern

tight-pattern ::=
```

```
        aliased-pattern
      | cons-pattern
      | embed-pattern
      | closed-pattern
```

```
closed-pattern  ::=
        variable-pattern
      | wildcard-pattern
      | literal-pattern
      | list-pattern
      | tuple-pattern
      | record-pattern
      | ( pattern )
```

(As for expressions, the distinctions tight- and closed- for patterns have no semantic significance, but merely serve to avoid grammatical ambiguities.)

annotated-pattern ::= pattern : type-descriptor

aliased-pattern ::= variable-pattern **as** tight-pattern

cons-pattern ::= closed-pattern **::** tight-pattern

embed-pattern ::= constructor [closed-pattern]

variable-pattern ::= local-variable

wildcard-pattern ::= _

literal-pattern ::= literal

list-pattern ::= **[** list-pattern-body **]**

list-pattern-body ::= [pattern { **,** pattern }*]

tuple-pattern ::= **(** tuple-pattern-body **)**

tuple-pattern-body ::= [pattern **,** pattern { **,** pattern }*]

record-pattern ::= **{** record-pattern-body **}**

record-pattern-body ::= [field-patterner { **,** field-patterner }*]

field-patterner ::= field-name [equals pattern]

Sample patterns:

```
(i, p) : Integer * Boolean
z as {re = x, im = y}
hd :: tail
Push {top, pop = rest}
embed Empty
x

_
#z
[0, x]
(c1 as (0, _), x)
{top, pop = rest}
```

Restriction. Like all polymorphic or type-ambiguous constructs, a pattern may only be used in a context where its type can be uniquely inferred.

Disambiguation. A single simple-name used as a pattern is an embed-pattern if it is a constructor of the type of the pattern. Otherwise, the simple-name is a variable-pattern.

Restriction. Each local-variable in a pattern must be a different simple-name, disregarding any local-variables introduced in expressions or type-descriptors contained in the pattern. (For example, `Line (z, z)` is not a lawful pattern, since `z` is repeated; but `n : {n : Nat | n < p}` is lawful: the second `n` is "shielded" by the type-comprehension in which it occurs.)

To define acceptance and acceptance binding (if any) for a value and a pattern, we introduce a number of auxiliary definitions.

The *accepting* local-variables of a pattern P are the collection of local-variables occurring in P, disregarding any local-variables introduced in expressions or type-descriptors contained in the P. For example, in pattern `u : {v : S | p v}`, `u` is an accepting local-variable, but `v` is not. (The latter is an accepting local-variable of pattern `v : S`, but not of the larger pattern.)

The *expressive descendants* of a pattern are a finite set of expressions having the syntactic form of patterns, as determined in the following three steps (of which the order of steps 1 and 2 is actually immaterial).

Step 1. From pattern P, form some *tame variant* P_t by replacing each field-patterner consisting of a single field-name F by the field-patterner $F = F$ and replacing each wildcard-pattern _ in P by a unique fresh simple-name, that is, any simple-name that does not already occur in the spec, directly or indirectly through an import. For example, assuming that the simple-name `v7944` is fresh, a tame variant of

```
s0 as _ :: s1 as (Push {top, pop = rest}) :: ss
```

is

```
s0 as v7944 :: s1 as (Push {top = top, pop = rest}) :: ss
```

Step 2. Next, from P_t, form a (tamed) *construed version* P_{tc} by replacing each constituent cons-pattern `H :: T` by the embed-pattern `Cons (H, T)`, where `Cons` denotes the constructor of the parameterized type `List`. For the example, the construed version is:

```
s0 as Cons (v7944,
            s1 as Cons (Push {top = top, pop = rest}, ss))
```

Step 3. Finally, from P_{tc}, form the set ED_P of *expressive descendants* of P, where expression E is an expressive descendant if E can be obtained by repeatedly replacing some constituent aliased-pattern `L as R` of P_{tc} by one of the two patterns L and R until no aliased-patterns remain, and then interpreting the result as an expression. For the example, the expressive descendants are the three expressions:

```
s0
Cons (v7944, s1)
Cons (v7944, Cons (Push {top = top, pop = rest}, ss))
```

An *accepting binding* of a pattern P for a value x in an environment C is some binding B of typed values to the accepting local-variables of the *tame* variant P_t, such that the value of each expressive descendant E in ED_P in the environment C extended with binding B, is the same typed value as x.

Acceptance and acceptance binding, if any, for a value x and a pattern P are then determined as follows. If there is no accepting binding of P for x, x is rejected. If an accepting binding exists, the value x is accepted by pattern P. There is a unique binding B among the accepting bindings in which the type of each assigned value is as "restricted" as possible in the subtype-supertype hierarchy without violating well-typedness constraints (in other words, there are no avoidable implicit coercions). The acceptance binding is then the binding B *projected on* the accepting local-variables of P.

For the example, the accepting local-variables of P_t are the six local-variables `s0`, `s1`, `ss`, `rest` and `v7944`. In general, they are the accepting local-variables of the original pattern together with any fresh simple-names used for taming. Let the value x being matched against the pattern be

```
Cons (Empty, Cons (Push {top = 200, pop = Empty}, Nil))
```

Under the accepting binding

```
s0 = Cons (Empty, Cons (Push {top = 200, pop = Empty}, Nil))
s1 = Cons (Push {top = 200, pop = Empty}, Nil)
ss = Nil
top = 200
rest = Empty
v7944 = Empty
```

the value of each E in ED_P amounts to the value x. Therefore, x is accepted by the original pattern, with acceptance binding

```
s0 = Cons (Empty, Cons (Push {top = 200, pop = Empty}, Nil))
s1 = Cons (Push {top = 200, pop = Empty}, Nil)
ss = Nil
top = 200
rest = Empty
```

obtained by "forgetting" the fresh simple-name `v7944`.

Appendix A. Metaslang Grammar

This appendix lists the grammar rules of the Metaslang specification language. These rules are identical to those of the Chapter on *Metaslang*. They are brought together here, without additional text, for easy reference.

Appendix B. Inbuilts and Base Libraries

This appendix provides a brief description of the types and operators that are either "inbuilt" or provided by the current base libraries. The base libraries are automatically imported by every user-defined spec. The title of each section of this appendix is the qualifier of the type- and op-names given therein. For example, the full name for op `++` described in Section "List" is `List.++`. However, for the *unary* operator – on integers, the qualifier is `Integer_`. Note, also, that inbuilts cannot be qualified.

For the sake of brevity, `infixl` is abbreviated below to `L` and `infixr` to `R`.

B.1. Inbuilts

Inbuilt Type

```
Boolean
```

Inbuilt Ops

Name	Fix-ity	Type	Description
`=`	R 20	`[a] a * a -> Boolean`	tests if the parameters are equal
`~=`	R 20	`[a] a * a -> Boolean`	tests if the parameters are unequal
`~`		`Boolean -> Boolean`	logical negation
`&&`	R 15	`Boolean * Boolean -> Boolean`	non-strict logical and
`||`	R 14	`Boolean * Boolean -> Boolean`	non-strict logical or
`=>`	R 13	`Boolean * Boolean -> Boolean`	non-strict logical implication
`<=>`	R 12	`Boolean * Boolean -> Boolean`	logical equivalence

Name	Fix-ity	Type	Description
<<	L 25	$\{x:A,...,y:B,...\}$ * $\{x:A,...,z:C,...\}$ -> $\{x:A,...,y:B,...,z:C,...\}$	see Section Applications under *record update*

B.2. Boolean

Ops

Name	Fix-ity	Type	Description
toString		Boolean -> String	converts logical value to string
show		Boolean -> String	same as toString
compare		Boolean * Boolean -> Comparison	compares two logical values

B.3. Integer

Types

```
type Integer

type NonZeroInteger = {i : Integer | i ~= 0}
```

Ops

Name	Fix-ity	Type	Description
`-`		`Integer -> Integer`	unary minus (has qualifier `Integer_!`)
`+`	L 25	`Integer * Integer -> Integer`	addition
`-`	L 25	`Integer * Integer -> Integer`	subtraction
`*`	L 27	`Integer * Integer -> Integer`	multiplication
`div`	L 26	`Integer * NonZeroInteger -> Integer`	division (truncates towards 0)
`rem`	L 26	`Integer * NonZeroInteger -> Integer`	remainder (x `rem` y = x - y * (x `div` y))
`<`	L 20	`Integer * Integer -> Boolean`	less-than
`<=`	L 20	`Integer * Integer -> Boolean`	less-than-or-equal
`>`	L 20	`Integer * Integer -> Boolean`	greater-than
`>=`	L 20	`Integer * Integer -> Boolean`	greater-than-or-equal
`abs`		`Integer -> Integer`	absolute value
`min`		`Integer * Integer -> Integer`	minimum
`max`		`Integer * Integer -> Integer`	maximum
`compare`		`Integer * Integer -> Comparison`	compares two integers
`toString`		`Integer -> String`	converts integer to string
`show`		`Integer -> String`	same as `toString`

Name	Fix-ity	Type	Description
`intToString`		`Integer -> String`	same as `toString`
`intConvertible`		`String -> Boolean`	tests if string is representation of integer
`stringToInt`		`(String | intConvertible) -> Integer`	converts "convertible" string to integer

B.4. Nat

Types

```
type Nat = {n : Integer | n >= 0}
type PosNat = {n : Nat | n > 0 }
```

Ops

Name	Fix-ity	Type	Description
`succ`		`Nat -> Nat`	successor
`pred`		`Nat -> Integer`	predecessor
`zero`		`Nat`	the natural number 0
`one`		`Nat`	the natural number 1
`two`		`Nat`	the natural number 2
`posNat?`		`Nat -> Boolean`	yields false for 0, true otherwise
`toString`		`Nat -> String`	converts natural number to string

Name	Fix-ity	Type	Description
`show`		`Nat -> String`	same as `toString`
`natToString`		`Nat -> String`	same as `toString`
`natConvertible`		`String -> Boolean`	tests if string is representation of natural number
`stringToNat`		`(String | natConvertible) -> Nat`	converts "convertible" string to natural number

B.5. Char

Type

```
type Char
```

Ops

Name	Fix-ity	Type	Description
`ord`		`Char -> Nat`	converts character to natural number
`chr`		`Nat -> Char`	converts natural number to character
`isAlpha`		`Char -> Boolean`	true for letters
`isNum`		`Char -> Boolean`	true for digits
`isAlphaNum`		`Char -> Boolean`	true for letters and digits
`isAscii`		`Char -> Boolean`	true for ASCII characters
`isLowerCase`		`Char -> Boolean`	true for lower-case letters

Name	Fix-ity	Type	Description
`isUpperCase`		`Char -> Boolean`	true for upper-case letters
`toUpperCase`		`Char -> Char`	converts to upper case
`toLowerCase`		`Char -> Char`	converts to lower case
`compare`		`Char * Char -> Comparison`	compares two character values
`toString`		`Char -> String`	converts character to string
`show`		`Char -> String`	same as `toString`

B.6. String

Type

```
type String
```

Ops

Name	Fix-ity	Type	Description
`explode`		`String -> List(Char)`	converts string to list of characters
`implode`		`List(Char) -> String`	converts list of characters to string
`length`		`String -> Nat`	length of a string
`leq`	`L 20`	`String * String -> Boolean`	lexicographic less-than-or-equal

Name	Fix-ity	Type	Description
`lt`	`L 20`	`String * String -> Boolean`	lexicographic less-than
`++`	`L 25`	`String * String -> String`	string concatenation
`^`	`L 25`	`String * String -> String`	same as `++`
`concat`		`String * String -> String`	prefix op for string concatenation
`concatList`		`List String -> String`	returns the concatenation of the list elements
`sub`		`String * Nat -> Char`	returns the nth character in a string, counting from 0
`substring`		`String * Nat * Nat -> String`	`substring(s, m, n)` returns the substring of s from position m through position n-1, counting from 0
`map`		`(Char -> Char) * String -> String`	returns the concatenation of the results of applying the function given as first parameter to each character of the string
`translate`		`(Char -> String) * String -> String`	returns the concatenation of the results of applying the function given as first parameter to each character of the string
`all`		`(Char -> Boolean) * String`	true if all characters in the string satisfy the predicate given as first parameter
`exists`		`(Char -> Boolean) * String`	true if some character in the string satisfies the predicate given as first parameter
`newline`		`String`	the string representing a line break
`toScreen`		`String -> ()`	prints the string on the terminal

Name	Fix-ity	Type	Description
`writeLine`		`String -> ()`	same with a newline appended
`compare`		`String * String -> Comparison`	compares two strings

B.7. List

Type

```
type List a = | Nil | Cons a * List a
```

Ops

Name	Fix-ity	Type	Description
`nil`		`[a] List a`	the empty list
`null`		`[a] List a -> Boolean`	true for empty lists
`length`		`List a -> Nat`	length of a list
`cons`		`[a] a * List a -> List a`	constructs a list consisting of a first element and a list tail
`insert`		`[a] a * List a -> List a`	same as `cons`
`hd`		`[a] List a -> a`	returns the first element of the list
`tl`		`[a] List a -> List a`	returns the list tail without the first element
`++`	L 25	`[a] List a * List a -> List a`	list concatenation

Name	Fix-ity	Type	Description
`concat`		`[a] List a * List a -> List a`	prefix op for list concatenation
`flatten`		`[a] List(List(a)) -> List a`	returns the concatenation of the list elements
`diff`		`[a] List a * List a -> List a`	list subtraction: `diff(`x`,`y`)` returns a list containing the elements of x that are not in y, preserving the order of the elements in x
`member`		`[a] a * List a -> Boolean`	list membership
`nth`		`[a] List a * Nat -> a`	`nth(`x`, `n`)` returns the element at position n of list x, counting from 0
`nthTail`		`[a] List a * Nat -> List a`	`nthTail(`x`, `n`)` returns the tail of list x, starting after position n, counting from 0
`sublist`		`[a] List a * Nat * Nat -> List a`	$\text{sublist}(x, m, n)$]] returns the tail of list x, from position m up to but not including n, counting from 0
`foldl`		`[a,b] (a*b -> b) -> b -> List a -> b`	`foldl` f e x successively applies function f to the elements of list x from left to right. The second argument to f is initially e and at each next step the result of the previous invocation of f
`foldr`		`[a,b] (a*b -> b) -> b -> List a -> b`	like `foldl`, but the elements of the list are processed from right to left
`map`		`[a,b] (a -> b) -> List a -> List b`	applies function to each element of a list and returns the list consisting of the results

Name	Fix-ity	Type	Description
mapPartial		`[a,b] (a -> Option b) -> List a -> List b`	like `map` but replacing each result `Some` *y* by *y* and deleting `None` results.
filter		`[a] (a -> Boolean) -> List a -> List a`	returns the list of elements satisfying the given predicate
rev		`[a] List a -> List a`	reverse list
all		`[a] (a -> Boolean) -> List a -> Boolean`	true if all elements of the list satisfy the predicate given as first parameter
exists		`[a] (a -> Boolean) -> List a -> Boolean`	true if some element of the list satisfies the predicate given as first parameter
find		`[a] (a -> Boolean) -> List a -> Option(a)`	returns `Some` *x* where *x* is the first element in the list (from left to right) for which the given predicate yields true; if no such element exists, `None` is returned
tabulate		`[a] Nat * (Nat -> a) -> List a`	`tabulate(`*n*`, `*f*`)` returns the list `[`*f*`(0), `*f*`(1), ... , `*f*`(`*n*`-1)]`
firstUpTo		`[a] (a -> Boolean) -> List a -> Option (a * List a)`	returns `Some(`*e*`, `*x*`)` where *e* is the first element in the list (from left to right) satisfying the given predicate and *x* the initial list segment preceding *e*; if no such element exists, `None` is returned
splitList		`[a] (a -> Boolean) -> List a -> Option (List a * a * List a)`	returns `Some(`*x*`, `*e*`, `*y*`)` where *e* is the first element in the list (from left to right) satisfying the given predicate, *x* the initial list segment preceding *e*, and *y* the list tail following *e*; if no such element exists, `None` is returned

Name	Fix-ity	Type	Description
locationOf		`[a] List a * List a -> Option (Nat * List a)`	`locationOf(`s`, `t`)` returns `Some(`n`, `x`)` where n is the first position in list t (counting from from left to right) where list s occurs as a contiguous sublist, and x the list tail segment following s in t; if s does not occur in t, `None` is returned
compare		`[a] (a * a -> Comparison) -> List a * List a -> Comparison`	compares two list using the comparison function given as first parameter
show		`[a] String -> List String -> String`	`show(`s`, `x`)` returns the element strings in x concatenated, with string s inserted between any two elements

B.8. Compare

Type

```
type Comparison = | Less | Equal | Greater
```

Ops

Name	Fix-ity	Type	Description

Name	Fix-ity	Type	Description
`compare`		`Comparison * Comparison -> Comparison`	compares comparison values
`show`		`Comparison -> String`	converts comparison value to string

B.9. Option

Type

```
type Option a = | Some a | None
```

Ops

Name	Fix-ity	Type	Description
`some`		`[a] a -> Option a`	op that constructs `Some x`
`none`		`[a] Option a`	op that constructs `None`
`some?`		`[a] Option a -> Boolean`	tests if the parameter is of the form `Some x`
`none?`		`[a] Option a -> Boolean`	tests if the parameter is `None`
`compare`		`[a] (a * a -> Comparison) -> Option a * Option a -> Comparison`	returns the result of the comparison of the two optional values, where `None` is less than `Some x` for all x; if both optional values are of the form `Some x`, the comparison function given as first parameter is used to compute the result

Name	Fix-ity	Type	Description
`mapOption`		`[a,b] (a -> b) -> Option a -> Option b`	applies the function given as first parameter to the optional value if it is `Some` x, otherwise `None` is returned
`show`		`[a,b] (a -> String) -> Option a -> String`	converts optional value to string; if the optional value is `Some` x, it uses the function given as first parameter to convert x to a string

B.10. Functions

Types

```
type Injective(a,b) = ((a -> b) | injective?)

type Surjective(a,b) = ((a -> b) | surjective?)

type Bijective(a,b) = ((a -> b) | bijective?)
```

Ops

Name	Fix-ity	Type	Description
`id`		`[a] a -> a`	identity function
`o`	L 24	`[a,b,c] (b -> c) * (a -> b) -> (a -> c)`	function composition
`injective?`		`[a,b] (a -> b) -> Boolean`	injectivity predicate; non-constructive

Name	Fix-ity	Type	Description
`surjective?`		`[a,b] (a -> b) -> Boolean`	surjectivity predicate; non-constructive
`bijective?`		`[a,b] (a -> b) -> Boolean`	bijectivity predicate; non-constructive
`inverse`		`[a,b] Bijective(a,b) -> Bijective(b,a)`	inverts bijective function; non-constructive