

Specware 4.0 User Manual

Specware 4.0 User Manual

Copyright © 2002 by Kestrel Development Corporation

Copyright © 2002 by Kestrel Technology LLC

All rights reserved

The name Specware® is a registered trademark of Kestrel Development Corporation

Table of Contents

1. Installing Specware	5
1.1. Contents of Distribution Package	5
1.2. System Requirements.....	5
1.2.1. Hardware.....	6
1.2.2. Operating system	6
1.3. Installation Instructions.....	6
1.4. Uninstalling.....	7
2. Getting Started	8
2.1. Starting Specware	8
2.2. Exiting Specware	8
3. Usage Model	9
3.1. Units	9
3.2. Interaction	9
4. Defining Units	11
4.1. Conceptual Model.....	11
4.1.1. Unit Identifiers	11
4.1.2. Unit Terms	11
4.2. Realization via the File System.....	11
4.2.1. The SWPATH Environment Variable	12
4.2.2. Single Unit in a File	12
4.2.3. Multiple Units in a File	13
4.3. Unit Definitions Are Managed Outside of Specware	14
5. Processing Units	16
5.1. Overview	16
5.2. Resolution of Unit Identifiers.....	17
5.2.1. SWPATH-Based Unit Identifier.....	17
5.2.2. Relative Unit Identifiers	19
5.3. Processing Commands	21
5.3.1. Processing a Unit	21
5.3.2. Processing a Multiple-Unit File.....	22

5.4. Unit-Specific Processing	22
5.4.1. Typechecking Specs	22
5.4.2. Proving Properties in Specs	23
5.4.2.1. The Proof Unit	24
5.4.2.2. Proof errors	24
5.4.2.3. Proof Log Files	25
5.4.2.4. Multiple Proofs	25
5.4.2.5. Interrupting Snark	25
5.4.2.6. The Prover Base Library	26
5.4.2.7. The Experimental Nature of the Prover	26
5.4.3. Generating Lisp Code	26
5.5. Auxiliary Commands	27
5.5.1. Displaying Unit Values	27
5.5.2. Inspecting and Clearing the Cache	28
5.5.3. Inspecting and Setting <code>SWPATH</code>	28
5.5.4. Generating Lisp Code	29
6. Lisp Code Generated from Specs	31
6.1. Translation of Specware Names to Lisp Names	31
6.2. Arity and Currying Normalization	32
6.3. Representation of Other Sorts	33
7. Debugging Generated Lisp Files.....	35
7.1. Tracing	35
7.2. Breaking	35
7.3. Timing	36
7.4. Interrupting	37
8. Emacs Usage in Specware	38
8.1. Specware Mode Commands	38
8.2. Specware Interaction Commands	39
8.3. Other Useful Emacs Commands	40

Chapter 1. Installing Specware

1.1. Contents of Distribution Package

The following programs and documents are included on the Specware 4.0 installation CD:

- `setup.exe` - This program in the root directory of the CD is the Specware 4.0 installer for Windows. It should be launched automatically when the CD is inserted into the CD-ROM drive.
- `README.txt` - This document contains important information about Specware 4.0, including what has changed from previous versions and how to give feedback on current and future features.
- `XEmacs` - This folder contains the XEmacs `setup.exe` installer for Windows. XEmacs is the environment within which Specware 4.0 is designed to run.
- `SNARK` - This folder contains the license agreement and modified source code for the theorem prover SNARK, which is built into Specware 4.0.

The following documentation is included with the distribution package:

- Specware 4.0 User Manual (this document) - The User Manual serves as a quick guide to basic usage and concepts of Specware. After reading this, you should feel comfortable with the mechanics of running and using Specware.
- Specware 4.0 Tutorial - The Tutorial will guide you through the process of specifying, refining and generating code in Specware. A comprehensive example provides step-by-step instructions for this development process.
- Specware 4.0 Language Manual - The Language Manual discusses the Metaslang specification language and gives the grammar rules and meaning for each Metaslang language construct.

1.2. System Requirements

1.2.1. Hardware

Specware has relatively modest system requirements for simple projects. Of course, as with any development tool, as your projects being developed become more complex, you may wish to work on a more powerful machine. For average use, however, the following basic hardware configuration is recommended:

- CPU: 250 Mhz
- RAM: 128 MB total, at least 64 MB free for applications
- Disk space: 15 MB for base system, 10-50 MB for user projects

1.2.2. Operating system

This version of Specware 4.0 has been tested to work with Windows XP, Windows 2000 and Windows NT 4.0.

1.3. Installation Instructions

1. Insert the installation CD into your CD-ROM drive; the Specware 4.0 `setup.exe` installer wizard will be automatically launched.
2. After accepting the license agreement, the installer will check whether XEmacs is installed on your machine. If it is, a message will appear displaying the found path to XEmacs. If it is not, you will be asked whether you would like to install it. It is **HIGHLY** recommended that you answer `yes`, as Specware 4.0 has been designed to run inside the XEmacs environment. Clicking `yes` will then launch the XEmacs installer wizard. Select the "Install from Local Directory" option for quickest installation, and click `Next` through each step in the wizard to accept the default

configuration. After XEmacs has been installed, you will return to the Specware installer wizard.

3. Select the directory where you would like Specware to be installed (the default is `C:\Program Files\Specware`, and click `Next` to complete the installation. A shortcut to Specware will be placed on your Desktop as well as in the Program Files folder in the Start menu. Documentation, libraries and example code will be placed in the installation directory you selected.

1.4. Uninstalling

To uninstall Specware 4.0, run the installer on the CD again and select the "Remove" option in the wizard, or use the `Add/Remove Programs` setting in the Control Panel.

If you need to reinstall Specware, please uninstall it first and then run the installer on the CD to install a clean version.

Chapter 2. Getting Started

Specware is a development environment that runs on top of Lisp. The following sections describe the Specware environment and the basic mechanisms for running Specware.

2.1. Starting Specware

To start Specware, double-click the `Specware 4.0` shortcut on your Desktop, or select `Specware 4.0` from the `Start Menu -> Program Files -> Specware` folder. When Specware is launched, a couple things happen: XEmacs is started and a Lisp image containing Specware is started inside an XEmacs buffer. All of the user interaction (see the next chapter) with Specware occurs at the Lisp prompt.

2.2. Exiting Specware

To exit Specware, type `:exit` at the Lisp prompt. A message will appear indicating that another process exists. Type `y` to confirm that you want to exit. This will kill Specware and you may then close the XEmacs window.

Chapter 3. Usage Model

3.1. Units

Simply put, the functionality provided by Specware consists in the capability to construct specs, morphisms, diagrams, code, proofs, and other entities. All these entities are collectively called *units*.

Some of the operations made available by Specware to construct units are fairly sophisticated. Examples are colimits, extraction of proof obligations, discharging of proof obligations by means of external theorem provers, and code generation.

The Metaslang language is the vehicle to construct units. The language has syntax to express all the unit-constructing operations that Specware provides. The user defines units in Metaslang, writing the definitions in `.sw` files (this file extension comes from the first and fifth letter of “Specware”).

Currently, the only way to construct units in Specware is by writing text in Metaslang. Future versions of Specware will include the ability to construct units by other means. For instance, instead of listing the nodes and edges of a diagram in text, it will be possible to draw the diagram on the screen.

3.2. Interaction

All the interaction between the user and Specware takes place through the Lisp shell. The `.sw` files that define units are edited outside of Specware, i.e. using XEmacs, Notepad or any other text editor of choice. These files are processed by Specware by giving suitable commands at the Lisp shell.

When `.sw` files are processed by Specware, progress and error messages are displayed in the XEmacs buffer containing the Lisp shell. In addition, the results of processing are saved into an internal cache that Specware maintains. Lastly, processing of certain kinds of units result in new files being created. For example, when Lisp code is

generated from a spec, the code is deposited into a `.lisp` file.

Specware also features auxiliary commands to display information about units, inspect and clear the internal cache, and inspect and change an environment variable that determines how unit identifiers are resolved to `.sw` files.

Chapter 4. Defining Units

4.1. Conceptual Model

A unit definition consists of a unit identifier and a unit term. The identifier identifies the unit and the term defines how the unit is constructed.

An application developed with Specware consists of a set of unit definitions, some of which may come from libraries. Units have unique identifiers within the application.

4.1.1. Unit Identifiers

A unit identifier is a finite, non-empty sequence of word symbols (word symbols are defined in the Metaslang grammar). The sequence of word symbols is essentially a “path” in a tree: the units comprising an application are organized in a tree.

This provides a convenient and mundane way to organize the units comprising an application. Libraries are subtrees of the whole tree. Parallel development of different parts of an application can be carried out in different subtrees that can be later put together without risk of naming clashes.

4.1.2. Unit Terms

A unit term is text written in the Metaslang language. Metaslang features various ways to construct specs, morphisms, and all the other kinds of units. For instance, it is possible to construct a spec by explicitly listing its sorts, ops, and axioms. It is also possible to construct a spec by applying the colimit operation to a diagram of specs and morphisms.

A unit term may reference other units. For instance, a spec constructed by extending another one references the spec being extended.

4.2. Realization via the File System

The conceptual model just described is realized by means of the file system of the underlying operating system. The file system has essentially a tree structure. The tree of units comprising a Specware application is mapped to subtrees of the file system; the word symbols comprising a path are mapped to file and directory names.

Future versions of Specware will have a more sophisticated UI that will realize the conceptual model directly. Users will graphically see the units organized in a tree and they will be able to add, remove, move, and edit them. The mapping to the file system may even be made totally transparent to the user.

4.2.1. The `SWPATH` Environment Variable

The mapping of the conceptual unit tree to the file system is defined by the environment variable `SWPATH`. Similarly to the “path” environment variable in operating systems, `SWPATH` is a string consisting of a colon-separated list of absolute directory paths in the file system. See Section 5.5.3 for information on how to inspect and set `SWPATH`.

Roughly speaking, the unit tree consists of all the units defined in `.sw` files under the directories listed in `SWPATH`. The identifier of each unit is its path from the directory in `SWPATH` under which the file defining the unit is: if the unit is under directory `dir`, its identifier is its absolute path in the file system “minus” the `dir` prefix. This approximate statement is made precise and illustrated by examples below.

4.2.2. Single Unit in a File

The simplest way to define a unit is to write its term into a `.sw` file in the subtree of one of the directories in `SWPATH`. The identifier of the unit is the name of the file, without `.sw`, prefixed by the path from the directory in `SWPATH` to the file.

For example, suppose that `SWPATH` includes the directory `\users\me\specware` (assuming the Windows operating system; currently, since `SWPATH` is colon-separated, it is not possible to use drive letters in directories; Specware automatically prefixes directories with `c:`; this restriction will be removed in future versions of Specware).

The user creates a file named `A.sw` immediately under the directory `c:\users\me\specware\one\two`, containing the following text:

```
spec
  sort X
endspec
```

The absolute path of the file in the file system is `c:\users\me\specware\one\two\A.sw`. The unit is a spec containing just a sort `X`. The identifier of the unit is `/one/two/A`. Note that the path components are separated by “/” (forward slash), even though the underlying file system uses “\” (backward slash). Unit identifiers are sequences of word symbols separated by “/”, regardless of the underlying operating system.

4.2.3. Multiple Units in a File

It is also possible to put multiple units inside a `.sw` file. The file must be in the subtree of one of the directories in `SWPATH`. Instead of just containing a unit term, the file contains one or more unit definitions. A unit definition consists of a word symbol, “=” (equal), and a unit term.

This case works almost exactly as if the file were replaced by a directory with the same name (without `.sw`) containing one `.sw` file for every unit defined therein. Each such file has the word symbol of the unit as name (plus `.sw`) and the term of the unit as content.

The only difference between the case of multiple units per file and the almost equivalent case where the file is replaced by a directory containing single-unit files, is that in the former case the last separator is not “/” but “#” (sharp). (This is reminiscent of the URI syntax, where subparts of a document are referred to using “#”.)

Suppose, as in the previous example, that `SWPATH` includes the directory `\users\me\specware`. The user creates a file named `three.sw` immediately under the directory `c:\users\me\specware\one\two`, containing the following text:

```
B = spec
```

```
    sort Y
endspec

C = spec
    sort Z
endspec
```

This file defines two specs, one containing just a sort `Y`, the other containing just a sort `Z`. The identifier of the first spec is `/one/two/three#B`, the identifier of the second spec is `one/two/three#C`.

As a particular instance of the case of multiple units per file, it is possible to have just one unit definition in the file. This is different from just having a unit term in a file. If the file contains a unit definition, then the word symbol at the left of “=” is part of the unit’s identifier, together with “#” and the file path (relative to the directory in `SWPATH`). If instead the file contains a unit term, then the identifier of the unit is the file path (relative to the directory in `SWPATH`), without any “#” and additional word symbol.

Despite the possibility of having one unit definition in a file, in this manual we use the term “multiple-unit file” to denote a file that contains one or more unit definitions. The term “single-unit file” is instead used to denote a file that only contains a unit term.

4.3. Unit Definitions Are Managed Outside of Specware

The `.sw` files are created, deleted, moved, and renamed by directly interacting with the file system of the underlying operating system.

The content of the `.sw` files can be edited with any desired text editor. A possibility is to use XEmacs, which is started when Specware is started and is used to interact with Specware. The XEmacs-Specware combo can be thought of as a (rather limited) Integrated Development Environment (IDE).

Note that unit definitions can be managed without running Specware at all. As described in the next chapter, Specware is used to process unit definitions. Future

versions of Specware will provide true IDE functionality: unit definitions will be also managed within Specware, and the mapping to the file system could be even made transparent to the user.

Chapter 5. Processing Units

5.1. Overview

Unit definitions are processed by Specware. The user instructs Specware to process units by supplying certain commands. Specware has access, via the Lisp runtime environment, to the underlying file system, so it can access the `.sw` files that define units. The environment variable `SWPATH` determines which `.sw` files are accessed by Specware to find unit definitions.

Processing a unit causes the recursive processing of the units referenced in that unit's term. For instance, if a spec A extends a spec B which in turns extends a spec C, then when A is processed, B and C are also processed. There must be no circularities in the chain of unit dependencies.

Processing causes progress and/or error messages to be displayed in the XEmacs buffer containing the Lisp shell (which is also where the user supplies commands to Specware). Progress messages inform the user that units were processed without error. Error messages provide information on the cause of errors, so that the user can fix them by editing the unit definitions. When an error occurs in the definition of some unit, Specware displays the `.sw` file containing the unit term in a separate XEmacs buffer, with the cursor positioned at the point containing the erroneous text.

The processing of certain kinds of units also results in the creation of new files as an additional side effect. For instance, Lisp programs are a kind of unit, constructed by the `generate` operation of Metaslang. A side effect of processing one such unit is that the resulting code is written into a `.lisp` file.

When Specware processes a unit, it saves the processing results into an internal cache, associating the results with the unit's identifier. By using this cache, Specware avoids unnecessary re-computations: it only re-processes the units whose files have changed since the last time they were processed. From the point of view of the final result, this caching mechanism is completely transparent to the user. However, it improves the performance and response time of the system.

5.2. Resolution of Unit Identifiers

Unit terms may reference units in the form of unit identifiers. A unit identifier is resolved to the unit's term, which is contained in a `.sw` file. Unit identifiers are either `SWPATH`-based or relative; these two kinds are syntactically distinguished from each other and are resolved in slightly different ways.

5.2.1. `SWPATH`-Based Unit Identifier

A `SWPATH`-based unit identifier starts with `"/`, followed by a `"/`-separated sequence of one or more path elements, where the last separator may be `#`. Examples are `/a/b/c`, `/d`, and `/e#f`.

Specware resolves a `SWPATH`-based unit identifier in the following steps:

1. If the unit identifier contains `#`, the `#` itself and the path element following it are removed, obtaining a `"/`-separated sequence of one or more path elements, preceded by `"/`. Otherwise, no removal takes place. Either way, the result of this first step is a `"/`-separated sequence of path elements preceded by `"/`.
2. The `"/` signs of the `"/`-separated sequence of path elements preceded by `"/`, resulting from the previous step, are turned into `\`; in addition, `.sw` is appended at the end. Otherwise, the `"/` signs are left unchanged and `.sw` is appended at the end. Either way, the result of this second step is a (partial) path in the file system.
3. The path resulting from the previous step is appended after the first directory of `SWPATH`. If the resulting absolute path denotes an existing file, that is the result of this third step. Otherwise, the same attempt is made with the second directory of `SWPATH` (if any). Attempts continue until a directory is found in `SWPATH` such that the absolute path obtained by concatenating the directory with the result of the previous step denotes an existing file; such a file is the result of this step. If no such directory is found, the unit identifier cannot be resolved and an error is signaled by Specware.
4. There are two alternative steps here, depending on whether or not the original unit identifier contains `#`.

- a. This is the case that the original unit identifier does not contain “#”. If the file resulting from the previous step is a single-unit file, i.e., it contains a unit term, that the final result of resolution. Otherwise, an error is signaled by Specware.
- b. This is the case that the original unit identifier contains “#”. The file resulting from the previous step must be a multiple-unit file, i.e., it must contain a sequence of one or more unit definitions. If this is not the case, the unit identifier cannot be resolved and an error is signaled by Specware. If that is the case, a unit definition is searched in the file, whose path elements (to the left of “=”) is the same as the path element that follows “#” in the original unit identifier. If no such unit definition is found, the unit identifier cannot be resolved and an error is signaled by Specware. If such a unit definition is found, its unit term (at the right of “=”) is the final result of resolution.

For example, consider a unit identifier `/a/b/c`. Since it does not contain “#”, the first step does not do anything. The result of the second step is `\a\b\c.sw`. Suppose that `SWPATH` is `c:\users\me\specware:\tmp`, that `c:\users\me\specware` does not contain any a subdirectory, and that `c:\tmp\a\b\c.sw` exists. The result of the third step is the file `c:\tmp\a\b\c.sw`. If such a file is a single-unit file, its content is the result of the fourth step.

As another example, consider a unit identifier `/e#f`. The result of the first step is `/e`. The result of the second step is `\e.sw`. Assuming `SWPATH` as before and that `c:\users\me\specware` contains a file `e.sw`, the file `c:\users\me\specware\e.sw` is the result of the third step. The file must be a multiple-unit file. Assuming that this is the case and that it contains a unit definition with path element `f`, its unit term is the result of the fourth step.

The directories in `SWPATH` are searched in order during the third step of resolution. So, in the last example, if the directory `c:\tmp` also contains a file `e.sw`, such a file is ignored. This features can be used, for example, to shadow selected library units that populate certain file system directories in `SWPATH`.

For example, suppose that the first directory in `SWPATH` is `c:\specware\libs` and that the directory `c:\specware\libs\data-structures` contains files `Sets.sw`, `Bags.sw`, `Lists.sw`, etc. defining specs of sets, bags, lists, etc. The unit identifier

`/data-structures/Sets` resolves to the content of the file `c:\specware\libs\data-structures\Sets.sw`. If the user wanted to experiment with a slightly different version of the spec for sets, it is sufficient to prepend another directory to `SWPATH`, e.g. `c:\shadow-lib`, and to create the slightly different version of the spec for sets in `c:\shadow-lib\data-structures\Sets.sw`. The same unit identifier `/data-structures/Sets` will now resolve to the new version.

5.2.2. Relative Unit Identifiers

A relative unit identifier is a “/”-separated sequence of one or more path elements, where the last separator can be “#”. Examples are `a/b/c`, `d`, and `e#f`. So, and relative unit identifiers can be distinguished by the presence or absence of “/” at the beginning.

The resolution of relative unit identifiers does not depend on `SWPATH`, but on the location of the file where the unit identifier occurs. There are two cases to consider: the unit identifier occurring in a single-unit file and the unit identifier occurring in a multiple-unit file.

Suppose that the relative unit identifier occurs in a single-unit file. Then Specware attempts to resolve the unit identifier in the following steps:

1. If the unit identifier contains “#”, the “#” itself and the path element following it are removed, obtaining a “/”-separated sequence of one or more path elements. Otherwise, no removal takes place. Either way, the result of this first step is a “/”-separated sequence of path elements.
2. The “/” signs of the “/”-separated sequence of path elements, resulting from the previous step, are turned into “\”; in addition, `.sw` is appended at the end. Otherwise, the “/” signs are left unchanged and `.sw` is appended at the end. Either way, the result of this second step is a (partial) path in the file system.
3. The path resulting from the previous step is appended after the absolute path of the directory of the file containing the relative unit identifier. If the resulting absolute path denotes an existing file, that is the result of this third step. Otherwise, the unit identifier cannot be resolved and an error is signaled by Specware.

4. There are two alternative steps here, depending on whether the original unit identifier contains “#” or not.
 - a. This is the case where the original unit identifier does not contain “#”. If the file resulting from the previous step is a single-unit file, i.e., it contains a unit term, that is the final result of resolution. Otherwise, an error is signaled by Specware.
 - b. This is the case that the original unit identifier contains “#”. The file resulting from the previous step must be a multiple-unit file, i.e., it must contain a sequence of one or more unit definitions. If this is not the case, the unit identifier cannot be resolved and an error is signaled by Specware. If that is the case, a unit definition is searched in the file, whose path element (at the left of “=”) is the same as the path element that follows “#” in the original unit identifier. If no such unit definition is found, the unit identifier cannot be resolved and an error is signaled by Specware. If instead such a unit definition is found, its unit term (to the right of “=”) is the final result of resolution.

So, resolution of a relative unit identifier occurring in a single-unit file is like resolution of a SWPATH-based unit identifier, except that the directory where the identifier occurs is used instead of SWPATH.

Suppose, instead, that the relative unit identifier occurs in a multiple-unit file. Then Specware attempts to resolve the unit identifier in the following steps:

1. If the relative unit identifier is a single path element, Specware attempts to find a unit definition with that path element inside the file where the unit identifier occurs. If such a unit definition is found, its unit term is the final result of resolution. Otherwise, the following steps are carried out:
2. If the unit identifier contains “#”, the “#” itself and the path element following it are removed, obtaining a “/”-separated sequence of one or more path elements. Otherwise, no removal takes place. Either way, the result of this first step is a “/”-separated sequence of path elements.
3. The “/” signs of the “/”-separated sequence of path elements, resulting from the previous step, are turned into “\”; in addition, .sw is appended at the end.

Otherwise, the “/” signs are left unchanged and `.sw` is appended at the end. Either way, the result of this second step is a (partial) path in the file system.

4. The path resulting from the previous step is appended after the absolute path of the directory of the file containing the relative unit identifier. If the resulting absolute path denotes an existing file, that is the result of this third step. Otherwise, the unit identifier cannot be resolved and an error is signaled by Specware.
5. There are two alternative steps here, depending on whether the original unit identifier contains “#” or not.
 - a. This is the case that the original unit identifier does not contain “#”. If the file resulting from the previous step is a single-unit file, i.e., it contains a unit term, that is the final result of resolution. Otherwise, an error is signaled by Specware.
 - b. This is the case that the original unit identifier contains “#”. The file resulting from the previous step must be a multiple-unit file, i.e., it must contain a sequence of one or more unit definitions. If this is not the case, the unit identifier cannot be resolved and an error is signaled by Specware. If that is the case, a unit definition is searched in the file, whose path element (at the left of “=”) is the same as the path element that follows “#” in the original unit identifier. If no such unit definition is found, the unit identifier cannot be resolved and an error is signaled by Specware. If instead such a unit definition is found, its unit term (to the right of “=”) is the final result of resolution.

So, resolution of a relative unit identifier occurring in a multiple-unit file is like resolution of a relative unit identifier occurring in a single-unit file, preceded by an attempt to locate the unit in the file where the identifier occurs, only in case such a unit identifier is a path element.

5.3. Processing Commands

The Specware command to process units is `:sw`. The user supplies this command in the XEmacs buffer of the Lisp shell, followed by an argument. The argument is a unit

identifier.

5.3.1. Processing a Unit

The command to process a unit is:

```
:sw <unit-identifier>
```

The unit identifier can be `SWPATH`-based or relative. If it is `SWPATH`-based, Specware attempts to resolve it as explained in Section 5.2.1. If it is relative, Specware attempts to resolve it as explained in Section 5.2.2, as if the unit identifier occurred in a single-unit file of the current Lisp directory. Either way, if resolution fails an error is signaled by Specware.

If instead resolution succeeds, Specware parses and evaluates the unit term that results from resolution. Parsing and evaluation carry out the computations to construct the unit; they are Specware’s “core” functionality. Parsing and evaluation implement the semantics of the Metaslang language.

If the unit term references other units, Specware recursively resolves the unit identifiers and parses and evaluates their unit terms.

5.3.2. Processing a Multiple-Unit File

The command to process a multiple-unit file is:

```
:sw <unit-identifier>
```

The unit identifier must not contain “#”. Specware attempts to resolve the unit identifier. If it is a relative unit identifier, it is resolved as if it occurred inside a single-unit file in the Lisp current directory. However, the file obtained at the third step must be a multi-unit file, and not a single-unit file. If it is indeed a multi-unit file, Specware parses and evaluates all the unit definitions inside the file.

5.4. Unit-Specific Processing

This section describes what happens when specific kinds of units are processed.

5.4.1. Typechecking Specs

The user can construct specs by explicitly listing the sorts, ops, and axioms comprising the spec, possibly after importing one or more spec. When these spec definitions are processed, Specware typechecks all the expressions that appear in the spec.

Typechecking means checking that the expressions are type-correct, according to the rules of the Metaslang language.

In general, only some of the ops and variables that appear in an expression have explicit type (i.e., sort) information. Typechecking also involves reconstructing the types of those ops and variables that lack explicit type information.

Typechecking is an integrated process that checks the type correctness of expressions while reconstructing the missing type information. This is done by deriving and solving type constraints from the expression. For instance, if it is known that an op f has sort $A \rightarrow B$ then the type of the variable x in the expression $f(x)$ must be A , and the type of the whole expression must be B .

If the missing type information cannot be uniquely reconstructed and/or if certain constraints are violated, Specware signals an error, indicating the textual position of the problematic expression.

Since the Metaslang type system features subsorts defined by arbitrary predicates, it is in general undecidable whether an expression involving subsorts is type-correct or not. When Specware processes a spec, it assumes that the type constraints arising from subsorts are satisfied, thus making typechecking decidable.

The proof obligations associated with a spec, which are extracted via the Metaslang `obligations` operation, include conjectures derived from the type constraints arising from subsorts. If all of these conjectures are discharged (using a theorem prover) then all the expressions in the spec are type-correct.

5.4.2. Proving Properties in Specs

Specware provides a mechanism for verifying the correctness of properties either specified in Metaslang specs or automatically generated as proof obligations arising from refinements or typechecking. Currently Specware comes packaged with the Snark first-order theorem prover. Interaction with Snark is through the proof unit described below.

5.4.2.1. The Proof Unit

The user invokes the Snark theorem prover by constructing a proof term. A typical proof term is of the form:

```
prove property in spec_term
  using hypothesis1 hypothesis2 ...
  options "(use-paramodulation t)
          (use-resolution nil)
          (use-hyperresolution t)"
```

If this term is in file `proof.sw` then issuing the command `:sw proof` will result in translating `hypothesis1`, `hypothesis2` and `property` to Snark and then invoking the Snark prover to try to prove `property` from `hypothesis1` and `hypothesis2` using the options in the `options` list. Note that the properties `property`, `hypothesis1`, and `hypothesis2` must all be names of claims (i.e. axioms, conjectures, or theorems) that appear in `spec_term`. Note also that `hypothesis1` and `hypothesis2` are required to appear earlier in `spec_term` than `property`. Most users will omit the `options` list. Additionally, the `using` part can be omitted as well. In this case all the properties that appear in `spec_term` prior to `property` will be used to prove `property`.

After Snark completes Specware will report on the success or failure of the Snark proof.

5.4.2.2. Proof errors

Specware will report an error if `property` does not occur in `spec_term` or if one of

the axioms do not occur in `spec_term` prior to `property`.

Snark will likely break into Lisp if the user inputs an incorrect option.

5.4.2.3. Proof Log Files

In the course of its execution Snark typically outputs a lot of information as well as a proof when it finds one. All this output can be overwhelming to the user, yet invaluable in understanding why his proofs succeeded or failed. To deal with all this output Specware redirects all the Snark output to log files. In our example above, which executed a proof in the file `proof.sw`, Specware will create a subdirectory called `snark` at the same level as `proof.sw`. In that directory a log file called `proof.log` will be created that contains all the Snark output.

5.4.2.4. Multiple Proofs

Just as there can be multiple units per file, there can be multiple proofs in single file. For example, in file `proof.sw` we could include more than one proof as follows:

```
p1  = prove prop1 using ax1 ax2
pla = prove prop1 using ax3
p2  = prove prop2
```

In this case Snark will be invoked three separate times, writing three different log files. In this case an additional subdirectory will be created under `snark`, called `proof`. The three log files will then be: `snark/proof/p1.log`, `snark/proof/pla.log`, and `snark/proof/p2.log`.

5.4.2.5. Interrupting Snark

As is the case with any first-order prover, Snark is likely to either loop forever or run for a longer time than the user can wait. The user can provide a time limit for Snark by using an appropriate option. However, there are likely to be times when the user wants to stop Snark in the middle of execution. He can do this by typing `Cntrl-C` `Cntrl-C` in

the `*common-lisp*` buffer. This will then interrupt Snark and place the user in the Lisp debugger. He can exit the debugger by issuing the `:pop` command. A log file will still be written that he can look at if he needs to.

5.4.2.6. The Prover Base Library

Specware has a base library that is implicitly imported by every spec. Unfortunately, the axioms in this library are not necessarily written to be useful by Snark. Instead of having Snark use these libraries we have created a separate base library for Snark. This library is located at `/Library/Base/ProverBase.sw`. The axioms in this spec are automatically sent to Snark as part of any proof.

5.4.2.7. The Experimental Nature of the Prover

Our experience with the current prover interface is very new and as such we are still very much experimenting with it and don't expect it to be perfect at this point in time. Many simple theorems will be provable. Some that the user thinks should be might not, and the user will be required to add further hypothesis and lemmas that may seem unnecessary. We are currently working on making this interface as robust and predictable as possible, and welcome any feedback the user can offer.

One area where the user can directly experiment is with the axioms that make up the `ProverBase`. The axioms that make up an effective prover library are best determined by an experimental evolutionary process. The user is welcome to play with the axioms in the `ProverBase`, by adding new ones or changing or deleting old ones. Keep in mind the goal is to have a single library that is useful for a wide range of proofs. Axioms that are specific to different proofs should be created in separate specs and imported where needed.

5.4.3. Generating Lisp Code

The user generates Lisp code from a spec by constructing a Lisp program unit via the

Metaslang `generate` operation. The optional string given as argument must be an absolute file path in the file system. The generated Lisp code is deposited into that file. The `.lisp` file extension can be omitted.

The string argument can be omitted if the spec term given as argument is a unit identifier. In this case, Specware deposits the generated code into the file `A.lisp`, where `A` is the rightmost path element comprising the unit identifier. The `A.lisp` file is deposited in a `lisp` subdirectory immediately under the directory of the file containing the unit term of the spec identified by the unit identifier given as argument to `generate`.

For example, suppose that the first directory in `SWPATH` is `\tmp` and that a spec is defined in the single-unit file `c:\tmp\one\two\A.sw`. Suppose that the Lisp program unit is defined by:

```
generate lisp /one/two/A
```

Then the code is deposited into the file `c:\tmp\one\two\lisp\A.lisp`.

As another example, suppose that `SWPATH` is as before and that a spec is defined in the multiple-unit file `c:\tmp\one\two\f.sw`, and that `B` is the path element associated with the spec. Suppose that the Lisp program unit is defined by:

```
generate lisp /one/two/f#B
```

Then the code is deposited into the file `c:\tmp\one\two\lisp\B.lisp`.

5.5. Auxiliary Commands

5.5.1. Displaying Unit Values

When a unit definition is processed, a unit value is produced. For example, a spec is essentially a set of sorts, ops, and axioms. A spec can be constructed by means of

various operations in the Metaslang language, but the final result is always a spec, i.e., a set of sorts, ops, and axioms.

The `:show` command is used to display the value of a unit. Like `:sw`, it is also supplied in the XEmacs buffer of the Lisp shell and it is followed by a one-line argument. The argument is a unit identifier.

The command to display a unit is:

```
:show <unit-identifier>
```

This command first processes the unit, exactly as if it were `:sw <unit-identifier>`. In addition, if processing does not yield errors, the value of the unit is displayed on screen. Of course, if the unit has been already processed via a `:sw` command and its unit term has not been changed after that, the `:show` command will access the results saved in Specware's internal cache.

The `:show` command serves to inspect the value of units constructed via the Metaslang operations. This is especially useful for beginning users as an aid to clarify the semantics of such operations.

5.5.2. Inspecting and Clearing the Cache

As already mentioned, Specware maintains an internal cache that associates processing results with unit identifiers.

A list of the units currently present in the cache is displayed on screen via the following command:

```
:list
```

The cache is cleared (i.e., re-initialized) via the following command:

```
:sw-init
```

Normally, there is no need to use these commands.

5.5.3. Inspecting and Setting `SWPATH`

The value of the `SWPATH` environment variable is displayed on screen via the following command:

```
:swpath
```

The value of the `SWPATH` environment variable is changed via the following command:

```
:swpath <string>
```

The `<string>` must be a semicolon-separated list of absolute directory paths of the underlying operating system, surrounded by double quotes. The “\” signs inside the directory paths must be escaped with “\”, because the string is processed by Lisp. For example, in order to set `SWPATH` to `c:\users\me` it is necessary to write `:swpath "c:\\users\\me"`.

Changes to `SWPATH` only apply to the currently running Specware session. If Specware is quit and then restarted, `SWPATH` loses the value assigned to it during the previous session, reverting to its default value.

5.5.4. Generating Lisp Code

Normally, Lisp code is generated by constructing a Lisp program unit via the Metaslang `generate` operation and by processing such unit via the `:sw` command.

Specware provides an additional command to accomplish the same result without actually creating an explicit Lisp program unit. The command is:

```
:swl <unit-identifier> <string>
```

The unit identifier must resolve to a spec, which is processed by Specware. If the spec is successfully processed, Specware generates Lisp code from it (according to the semantics of `generate`) and deposits the resulting code into the file whose path is given by the string.

The second argument to `:sw1`, i.e. the string, is optional. If it is not given, a file name is inferred as explained in Section 5.4.3 for `generate`.

Chapter 6. Lisp Code Generated from Specs

The translation of executable specs to Lisp code is straightforward for the most part as Lisp is a higher-order functional language. Functional expressions go to lambda expressions and most Specware sorts are implemented as list lists and vectors apart from the strings, numbers, characters and booleans which are implemented by the corresponding list datatypes. This guide is meant primarily to help the user in calling and debugging the functions generated from a spec, so we concentrate on the translation of op names to Lisp names and the implementation of sorts. The implementation details of procedural constructs such as pattern-matching are omitted. The interested user is free to examine the list code itself, which is simple but verbose for pattern-matching constructs.

6.1. Translation of Specware Names to List Names

Specware ops are implemented using list defuns if they are functions, defparameter otherwise. Their names are uppercased and put in the package with the same name as the qualifier, or `SW-SPEC` if unqualified. However, if the name is that of a built-in List symbol, the name is prepended with the character "!" and not uppercased. If the qualifier of the op is the same as a built-in list package then `-SPEC` is appended to the spec name to get the package name. For example, the list code for the spec:

```
A qualifying spec
  def two: Nat = 2
  def add1(x:Nat): Nat = x + 1
endspec
```

is

```
(DEFPACKAGE "A")
(IN-PACKAGE "A")
```

```
(DEFPARAMETER TWO 2)
(DEFUN ADD1 (X) (NAT-SPEC::|!+| X 1))
```

6.2. Arity and Currying Normalization

All Specware functions are unary. Multiple argument functions are modeled using either functions with product domains, or curried functions. For efficiency we wish to exploit Common Lisp's support of n-ary functions. Arity normalization aims to minimize unpacking and packing of products when passing arguments to functions with product domains, and currying normalization aims to minimize closure creation when calling curried functions. The saving is particularly important for recursive functions where there is saving at each recursive call, and in addition, currying normalization may permit the Common Lisp compiler to do tail recursion optimization. The naming scheme does not require knowledge of the definition of a function when generating calls to the function.

For each function whose argument is a product, two entry points are created: an n-ary function whose name is derived from the op as described above, and a unary function whose name has "-1" appended. E.g. for

```
op min : Integer * Integer -> Integer
```

there are two lisp functions #'MIN and #'MIN-1. A call with an explicit product is translated to the n-ary version, otherwise the unary version is used. For example, `min(1,2)` translates to `(MIN 1 2)`, and `foldr min inf 1` translates to `(FOLDR-1-1-1 #'MIN-1 INF L)`. When generating lisp for a definition, the form is examined to see whether the definition is naturally n-ary. If it is, then the primary definition is n-ary and the unary function is defined in terms of the n-ary function, otherwise the situation is reversed. For example, given the definition

```
def min(x,y) = if x <= y then x else y
```

we get the two Common Lisp definitions:


```
(DEFUN MIN (X Y) (if (<= x y) x y))
(DEFIN MIN-1 (X) (MIN (CAR X) (CDR X)))
```

and given the definition

```
def multFG(x: Nat * Nat) = (F x) * (G x)
```

we get the two Common Lisp definitions:

```
(DEFUN MULTFG (X Y) (PAIR1-1 (CONS X Y)))
(DEFUN MULTFG-1 (X) (* (F-1 X) (G-1 X)))
```

For each curried function (i.e. for each function whose codomain is a function) there is an additional uncurried version of the function with "-1" added n times to the name where n is the number of curried arguments. E.g. for

```
op foldr: fa(key,a,b) (a * b -> b) -> b -> map(key,a) -> b
```

there are two lisp functions #FOLDR and #FOLDR-1-1-1.

As with arity normalization, the definition of a curried function is examined to see whether it should be used to generate the curried or the uncurried version, with the other being defined in terms of this primary version.

As well as producing more efficient code, the currying normalization makes code easier to debug using the Common Lisp trace facility. For example if a function has a call of the form `foldr x y z`, this call is implemented as `(FOLDR-1-1-1 x y z)`, so you can trace `FOLDR-1-1-1` to find out how it is being called and what it is returning.

6.3. Representation of Other Sorts

`Character` and `String` sorts are represented as Lisp characters and strings, `Nat` and `Integer` as Lisp integers, lists are represented using Lisp lists, and `Boolean` `true` and `false` by the symbols `T` and `NIL`.

Sums are represented as the cons of the constructor name in keyword package and the fields of the constructor.

Binary products are implemented as cons cells (except for function arguments which are described in the previous section): `CONS` to construct and `CAR` and `CDR` to access the first and second fields. Non-binary products are implemented as vectors: constructed using `VECTOR` and the *i*th element accessed by `(SVREF x i-1)`.

Records are implemented the same as products with the order of the fields being alphabetic in the field names.

Restrictions and comprehensions are implemented using their supersort.

A quotient is represented as as a vector of three elements: the quotient tag (which is the value of the lisp variable `SLANG-BUILT-IN: :QUOTIENT-TAG`), the representation of the quotient relation, and the actual value in the underlying sort.

Chapter 7. Debugging Generated Lisp Files

7.1. Tracing

If you need to debug your application, there are a number of useful lisp facilities you should be aware of. The simplest trick is to trace some functions you care about to see what they are doing.

```
(trace foo)
  This will display the arguments to foo each time it is
  called, and will display the results each time it returns.
```

```
(untrace foo)
  This will turn off any tracing on foo.
```

7.2. Breaking

If you need a more detail view of runtime behavior, you might want to BREAK some functions you care about.

```
(trace (foo :break-all t))

  This will invoke the debugger each time foo is called,
  and upon each exit from foo.
```

Once you arrive in the debugger, the following commands are most useful:

:down <n>	Move to a deeper frame	<n> is optional
:up <n>	Move to a higher frame	<n> is optional
:zoom <n>	Display n frames	<n> is optional
:pri	Enter a dialog that lets you set printer	

Chapter 7. Debugging Generated Lisp Files

control variables. For example, setting depth to 5 and length to 10 will let you see the top level structure of expressions, while suppressing deep expressions and the tails of long expressions. Note that you'll need to specify values for many contexts, but just hitting return leaves a value unchanged. You will likely want to modify the trace, debugger, and current values.

(pprint *) Pretty print the expression for the current frame. Note that this only works immediately after arriving at a frame, e.g. via :down 0 if necessary.

:cont Continue as if nothing happened.
:restart Resume execution at this frame.
:reset Return to lisp top level. (E.g., bail out to try again.)
:exit Exit from lisp to operating system - ends session.

:help Online documentation.

(misc ...) The debugger is is a read/eval/print loop, so arbitrary lisp forms will be evaluated (in the current dynamic context).

(untrace foo) Stop entering debugger when foo is called.
Note that you may still enter the debugger for each exit from calls to foo already recursively in progress.

7.3. Timing

If you are curious about the overall performance of your application, the `TIME` macro will provide some quick information:

```
(time (foo nil))  
  This will report the time and space used by foo, e.g.:  
  
USER(1): (time (list 1 2 3))  
; cpu time (non-gc) 0 msec user, 0 msec system  
; cpu time (gc)      0 msec user, 0 msec system  
; cpu time (total)  0 msec user, 0 msec system  
; real time   231 msec  
; space allocation:  
; 6 cons cells, 0 symbols, 0 other bytes, 0 static bytes  
(1 2 3)
```

Note that `TIME` is transparent, i.e., it returns whatever its argument would return, including multiple values, etc., so it is safe to intersperse it nearly anywhere.

Common Lisp has more facilities for rolling your own timers: see the generic Common Lisp documentation, or contact Kestrel Technologies.

7.4. Interrupting

Finally, note that a useful trick in lisp is to start your application, e.g. `(foo nil)`, then at an appropriate time hit control-C. This will interrupt your application and put you into the debugger. From there you can enter the command `:zoom` to see the top of the stack. That can often be quite revealing.

Chapter 8. Emacs Usage in Specware

There are a number of emacs commands for Specware usage. We also list some useful general-purpose commands. There is a Specware mode for editing specware files.

8.1. Specware Mode Commands

These are the commands available when editing `.sw` files. Some of them are available in the Specware menu in the menubar.

Key	Name	Description
tab	sw:indent-line	Indents line based on context.
line-feed	newline-and-indent	Same as return followed by tab.
M-tab	sw:back-to-outer-indent	Indents line for outer expression.
M-C-q	sw:indent-sexp	Indents parenthesized expression following point.
M-C-\	sw:indent-region	Indents region.
M-*	sw:switch-to-lisp	Switch to <code>*common-lisp*</code> buffer.
M-.	sw:meta-point	Prompts for op or sort name (with default based on symbol at point), and goes to its definition. Searches only the units that have been loaded. First searches for definitions visible from the current unit.
M-,	sw:continue-meta-point	If previous meta-point command returned more than one definition, go to the next definition. Can be repeated.

Key	Name	Description
M-	sw:electric-pipe	Adds the skeleton of a new case to a case statement, properly indented.
C-c ;	sw:comment-region	Comment out the region. With a negative argument uncomments the region.
C-c p	sw:process-current-file	Does :sw on current file. Also available in dired where it applies to the file on the current line.
C-c C-p	sw:process-unit	Prompts for unitId to process. Defaults to unitId for the current file.
C-c C-p	cd-current-directory	Does a :cd in the *common-lisp* buffer to the directory of the current file. Also available in dired.

8.2. Specware Interaction Commands

These are the commands available in the *common-lisp* buffer.

Key	Name	Description
M-*	sw:switch-to-lisp	When already in the *common-lisp* buffer, goes to the previous source file buffer.
M-.	sw:meta-point	Prompts for op or sort name (with default based on symbol at point), and goes to its definition. Searches only the units that have been loaded.
M-p	fi:pop-input	Gets previous input.

Key	Name	Description
M-r	fi:re-search-backward-input	Gets previous input matching regular expression (prompted for). Typically you can get a previous input by typing in a small substring. Repeating the command without changing the expression will find earlier matches.

8.3. Other Useful Emacs Commands

These are a few general Emacs commands which are useful when using Specware. Commands with no key sequences are executed using M-x name.

Key	Name	Description
C-h	help	Help options. C-h m gives help about the current mode. Also see Help menu in menubar.
M-/	dabbrev-expand	Does symbol completion based on nearby words in buffer. Repeated key presses find additional completions.
C-sh-middle	mode-motion-copy	Copies the (highlighted) identifier or expression under the mouse to point.
	igrep	Grep for string in files. Brings up buffer with matching lines. Mouse middle on a line to go to it. (In Windows requires installation of Cygwin.)
	igrep-find	Like igrep but searches in all subdirectories.
	fgrep	Like igrep except uses fgrep to search.
	fgrep-find	Like igrep-find except uses fgrep to search. all subdirectories.

Key	Name	Description
C-x C-f	find-file	Prompts for file to edit. tab does filename completion.
C-x d	dired	Prompts for directory to edit. Note that commands sw:process-current-file and cd-current-directory described in the previous section are available in dired mode.
	viper-mode	Vi emulation mode for people who like to edit using vi commands. Documentation is available under C-h i.

